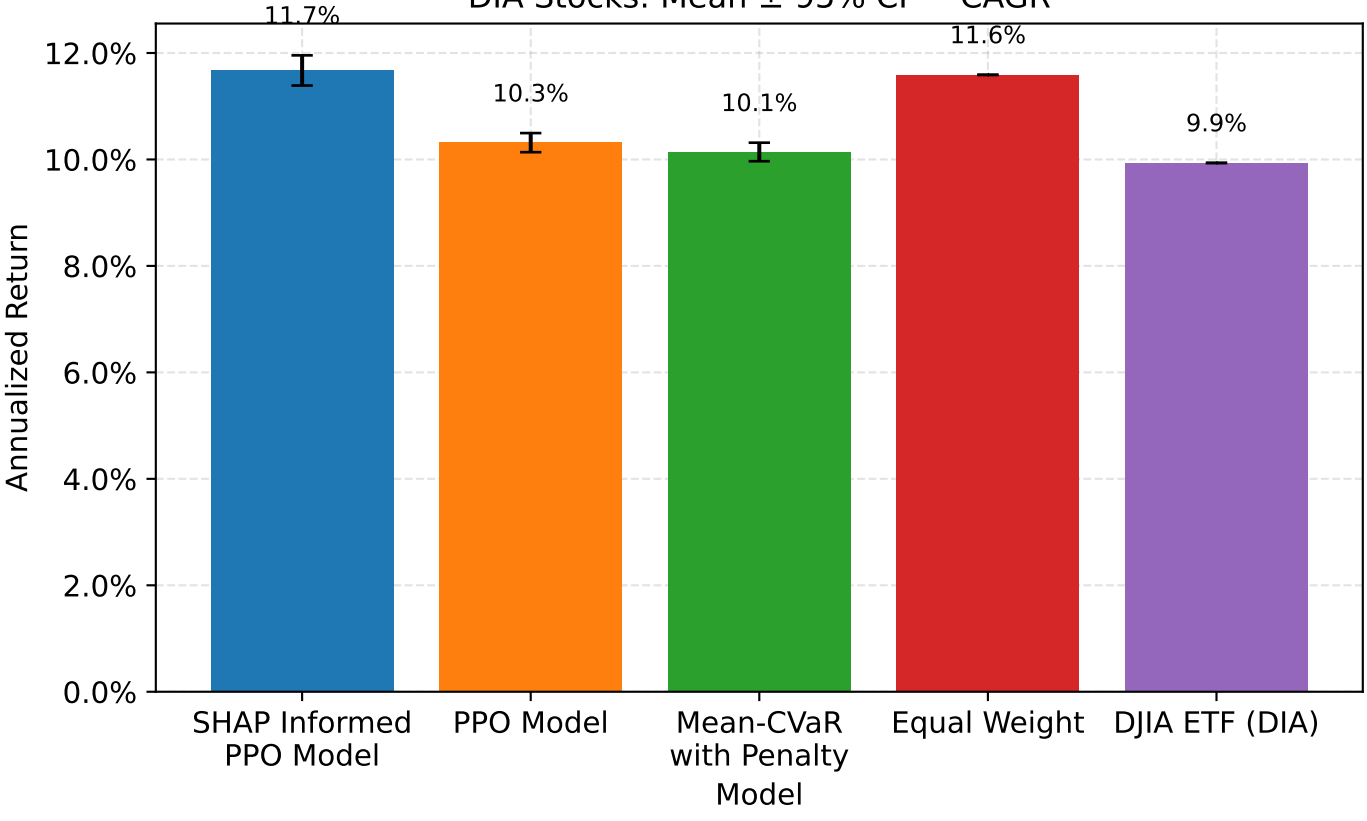

Bars show mean across 30 iterations; error bars are 95% normal-approximate CIs.

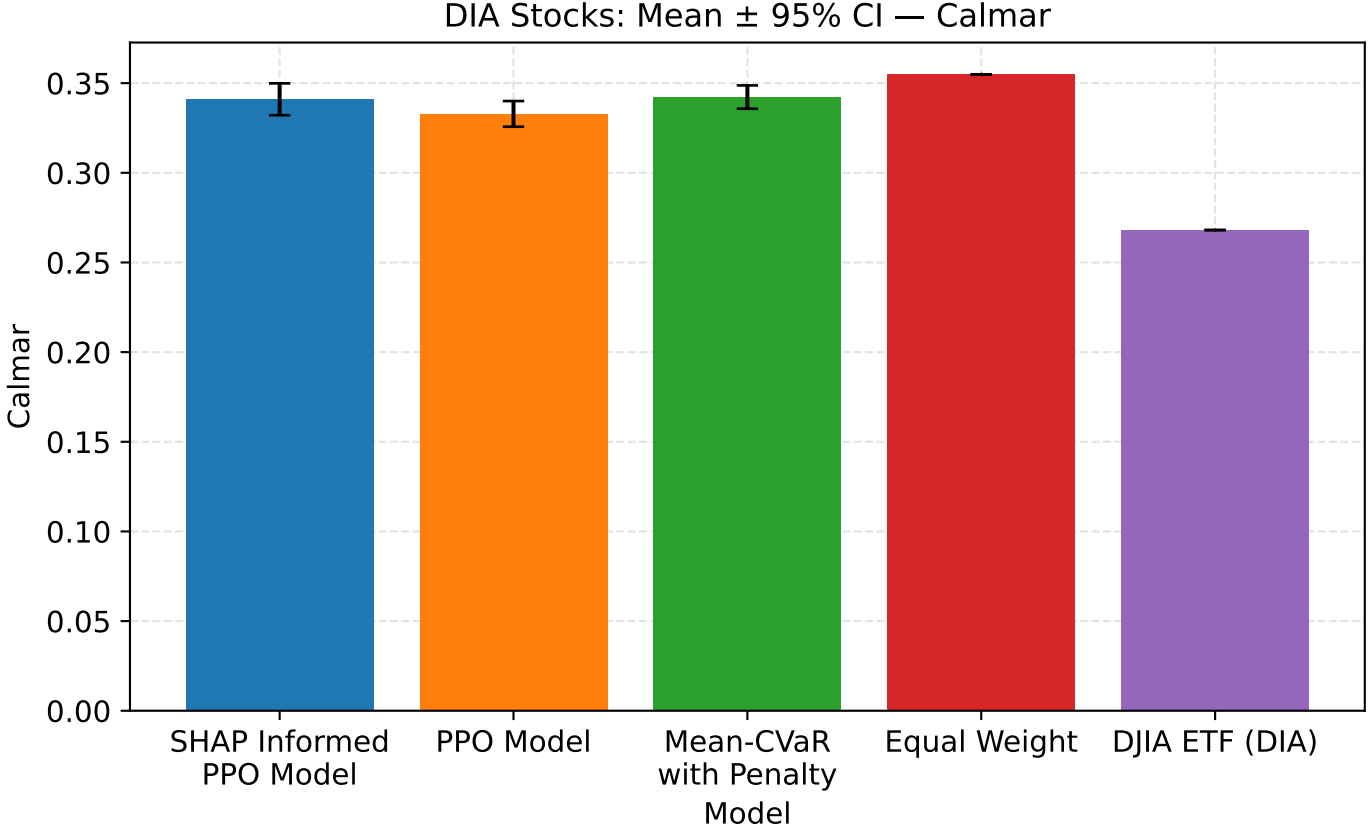

Bars show mean across 30 iterations; error bars are 95% normal-approximate CIs.

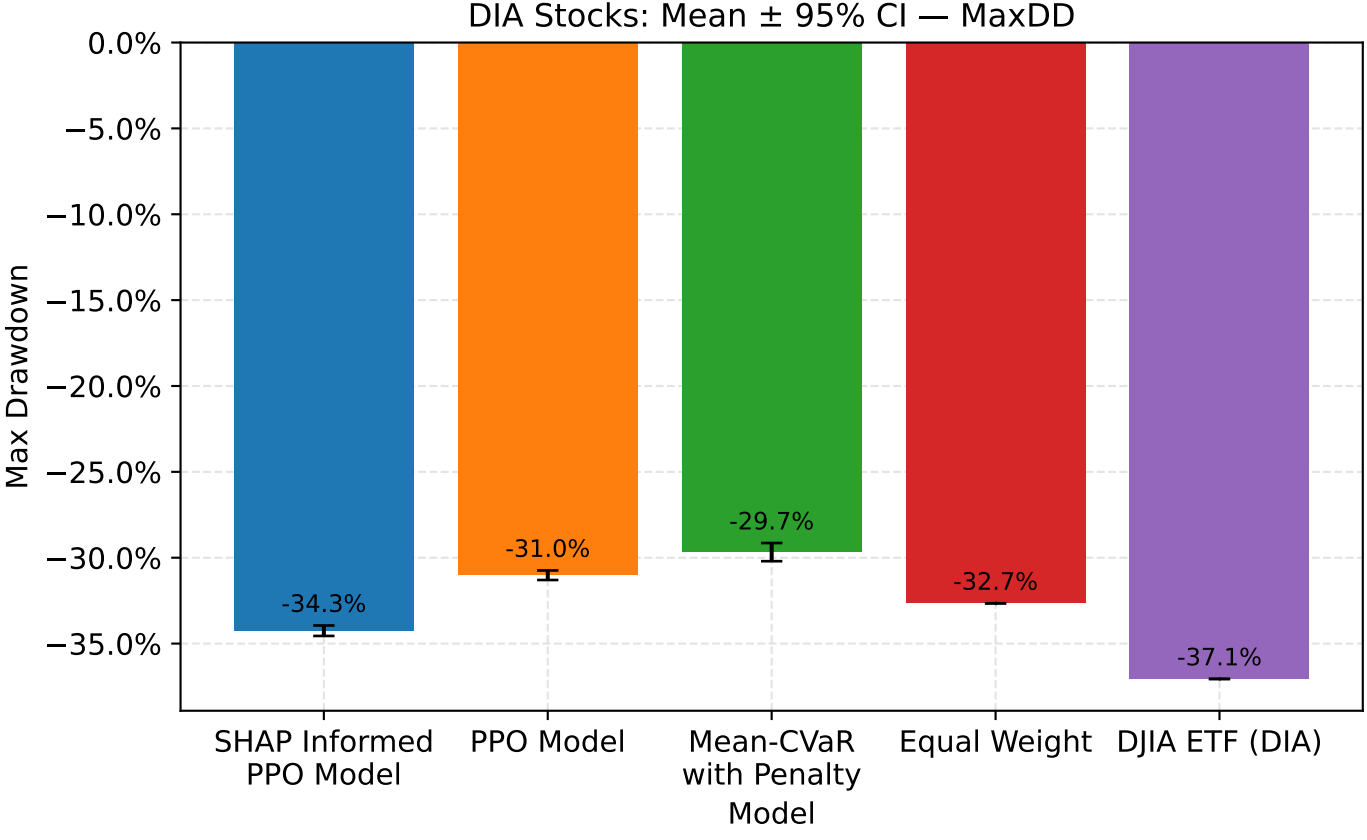

Bars show mean across 30 iterations; error bars are 95% normal-approximate CIs.

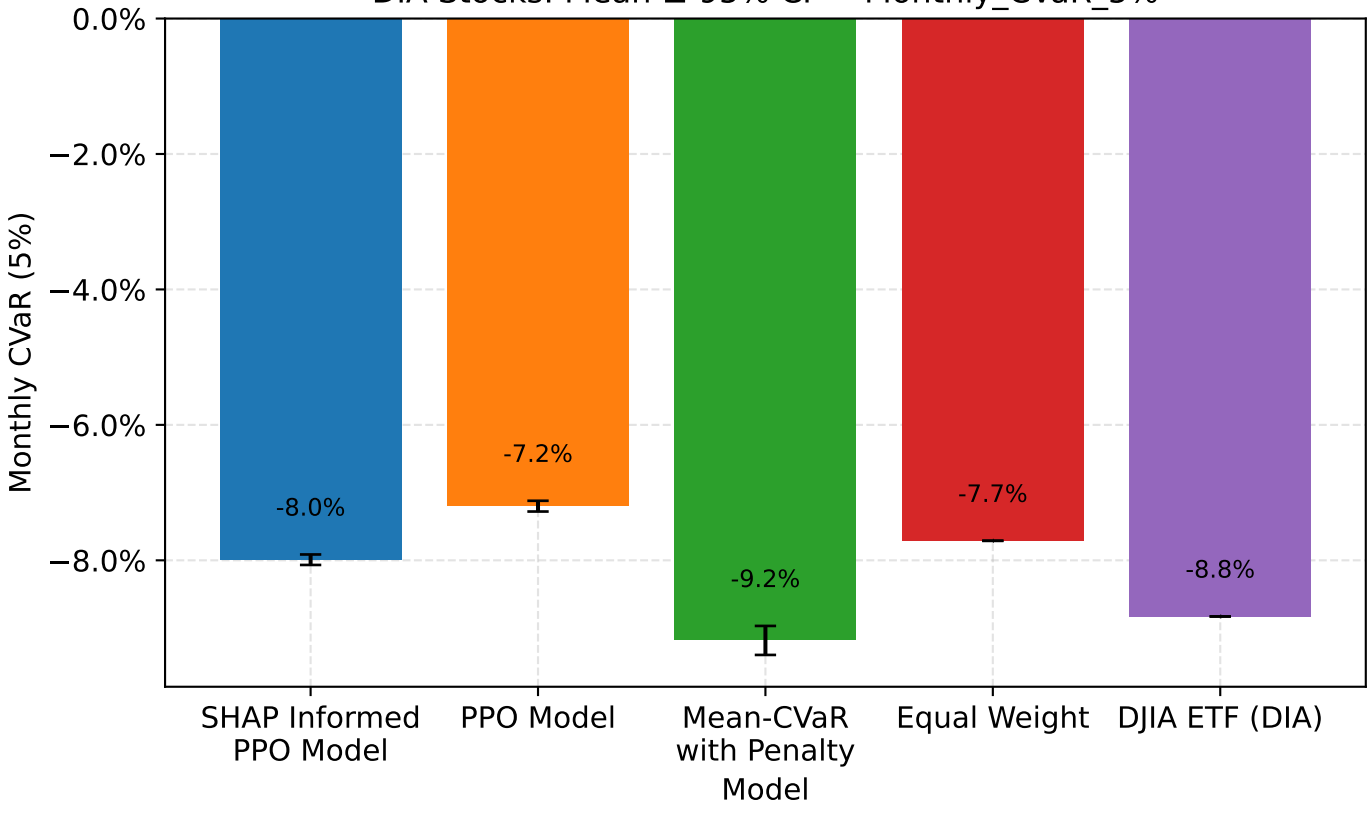

Bars show mean across 30 iterations; error bars are 95% normal-approximate CIs.

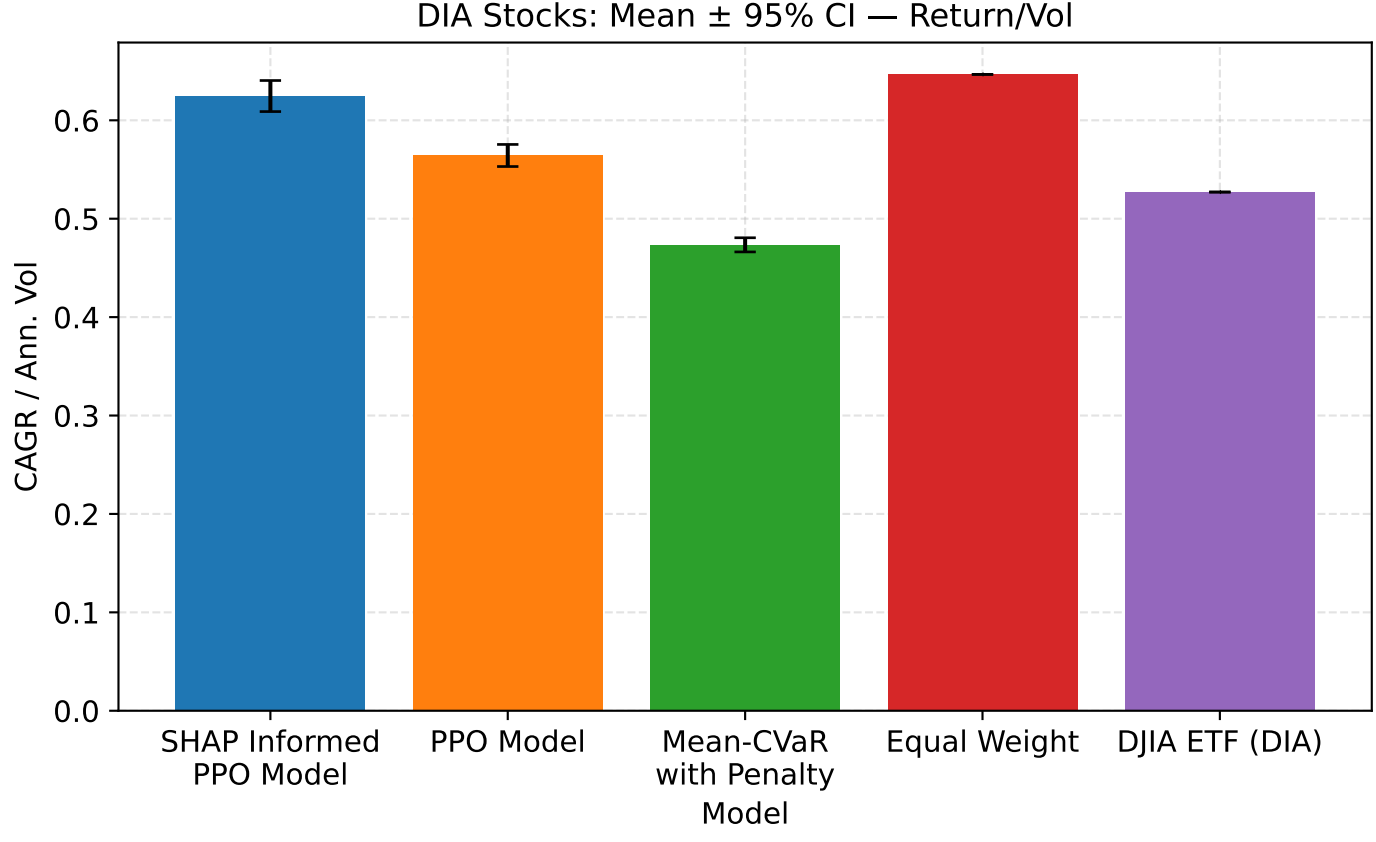

Bars show mean across 30 iterations; error bars are 95% normal-approximate CIs.

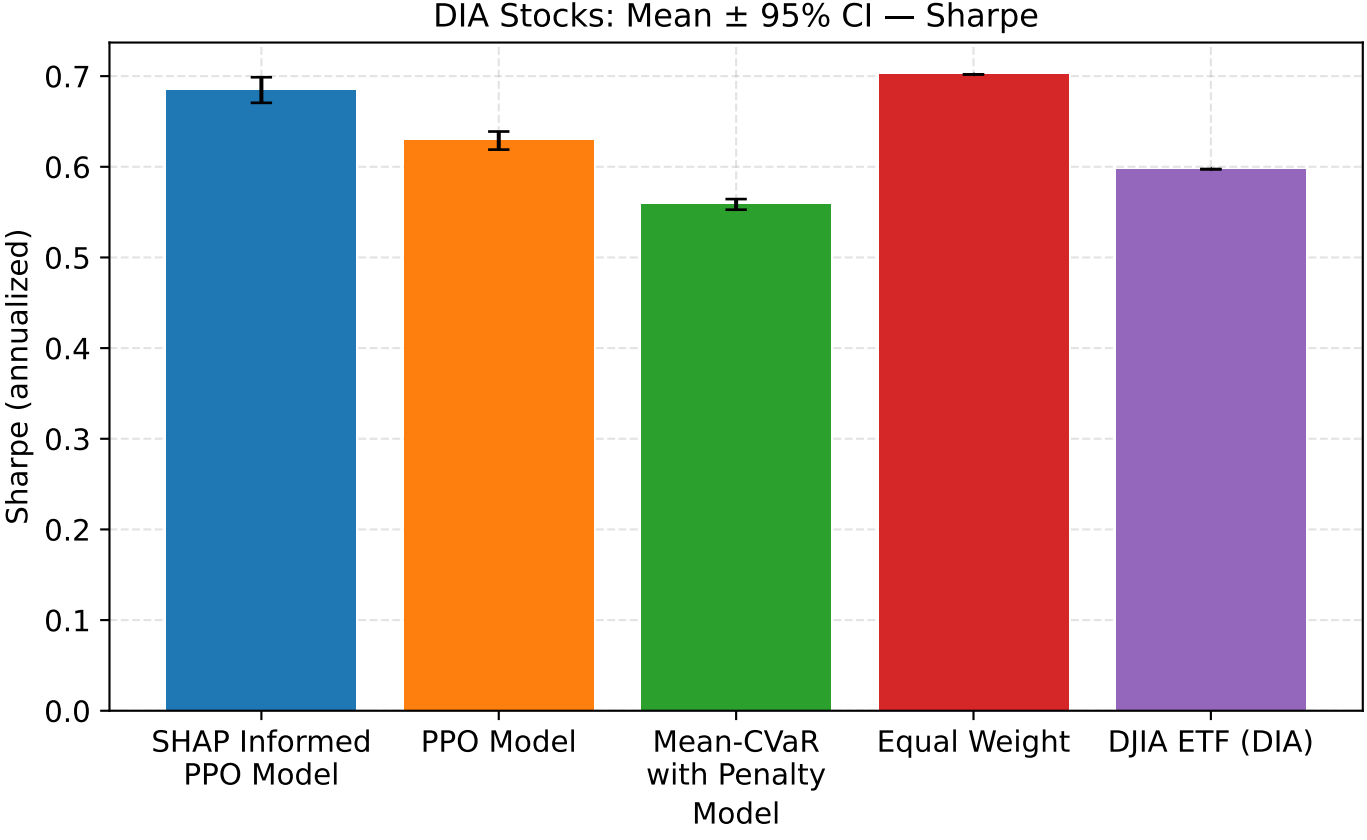

Bars show mean across 30 iterations; error bars are 95% normal-approximate CIs.

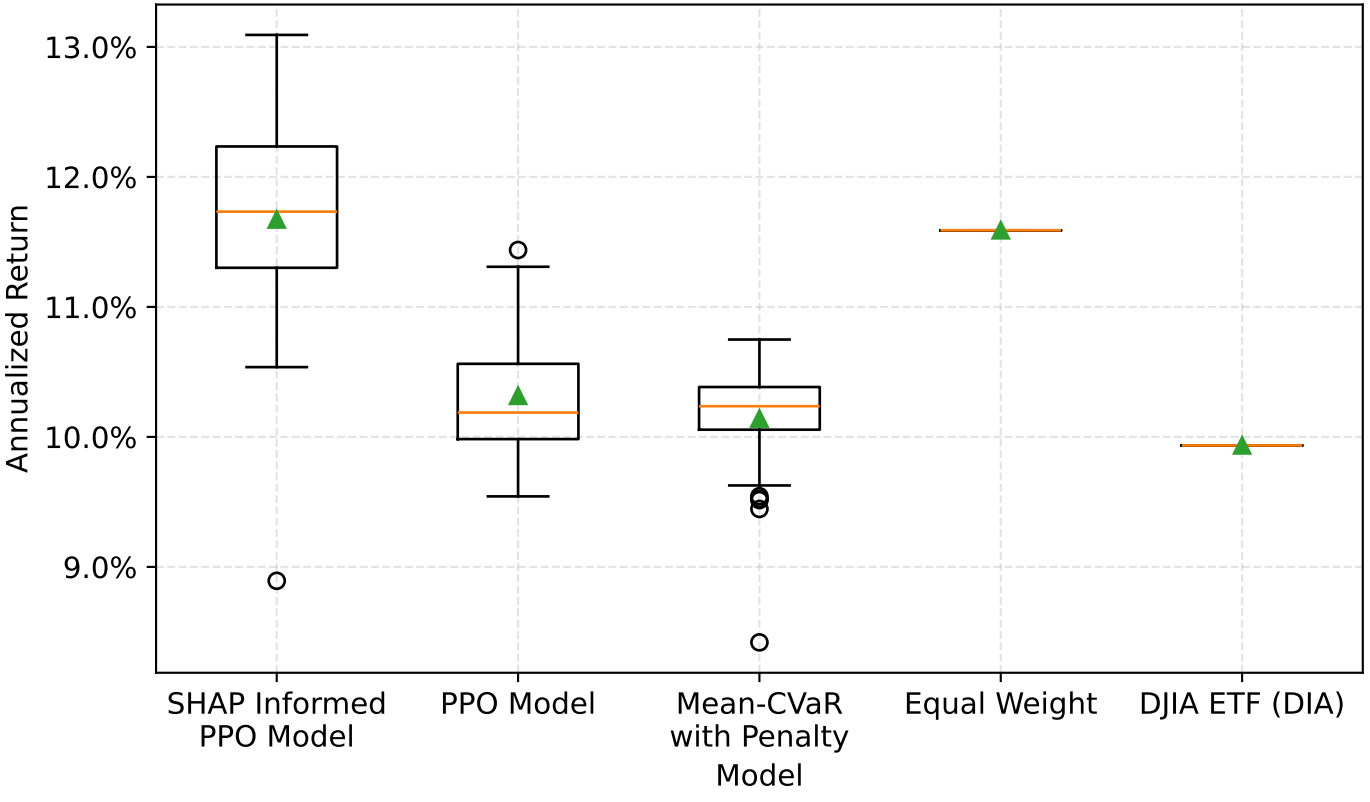

Boxes: IQR; center line: median; diamond = mean; whiskers: non-outlier range.

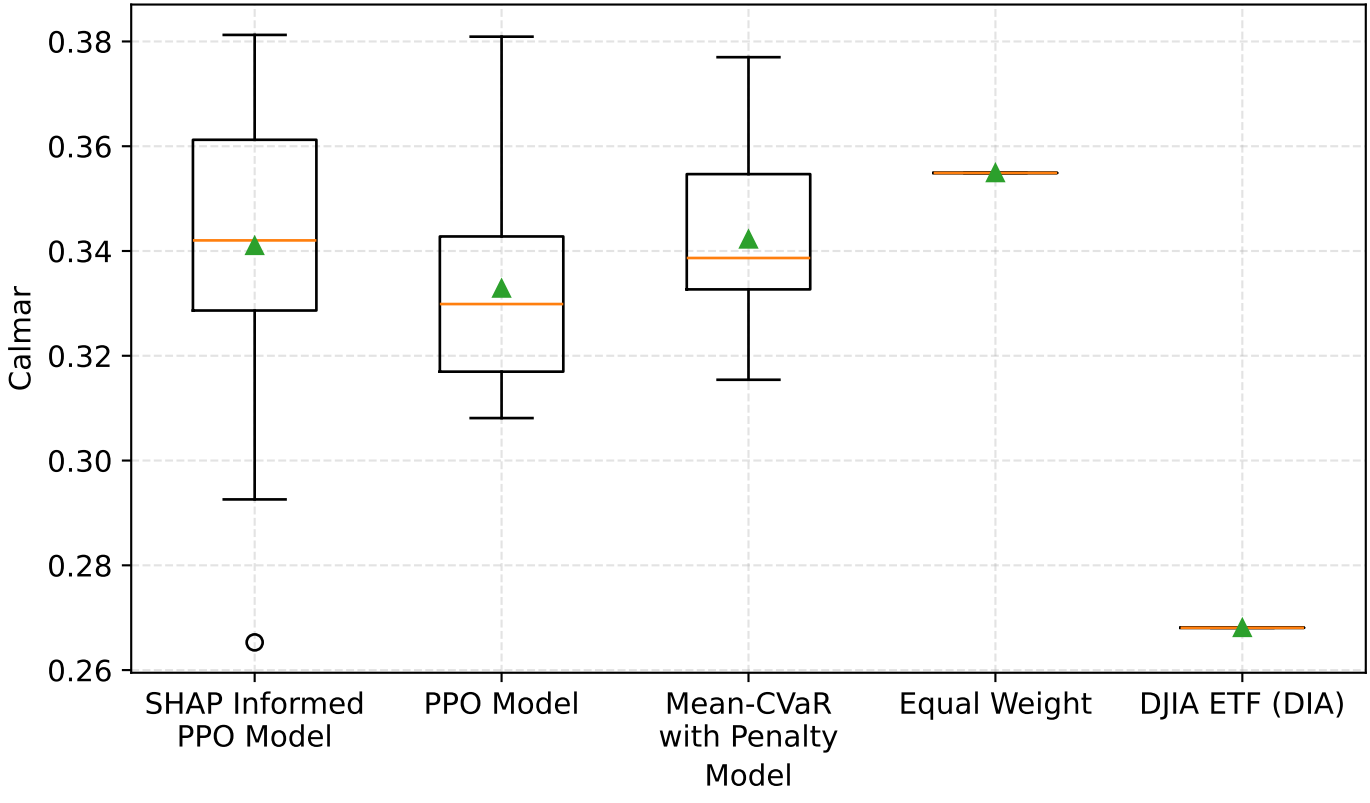

Boxes: IQR; center line: median; diamond = mean; whiskers: non-outlier range.

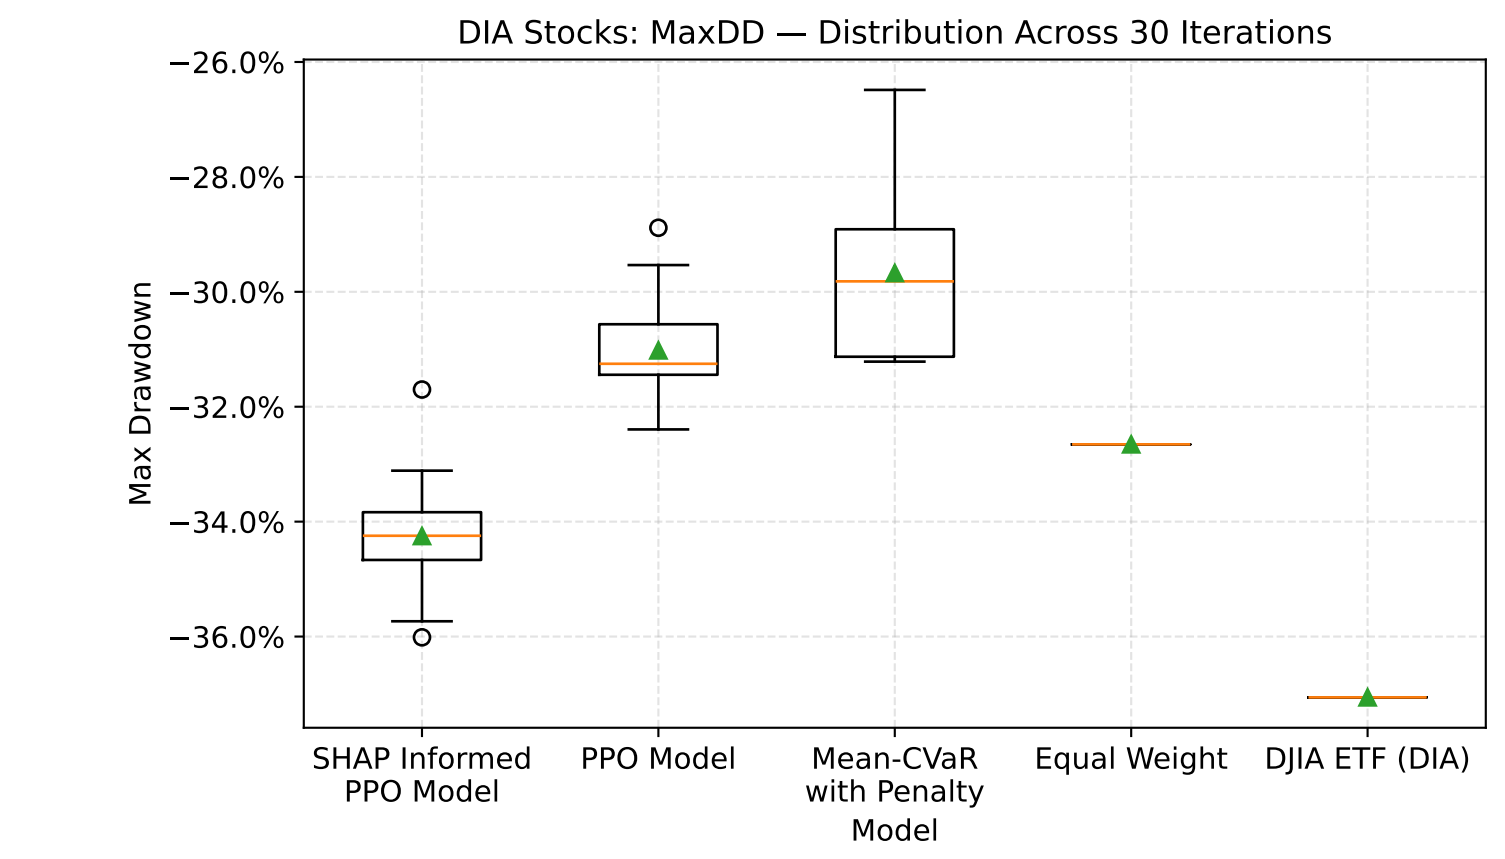

Boxes: IQR; center line: median; diamond = mean; whiskers: non-outlier range.

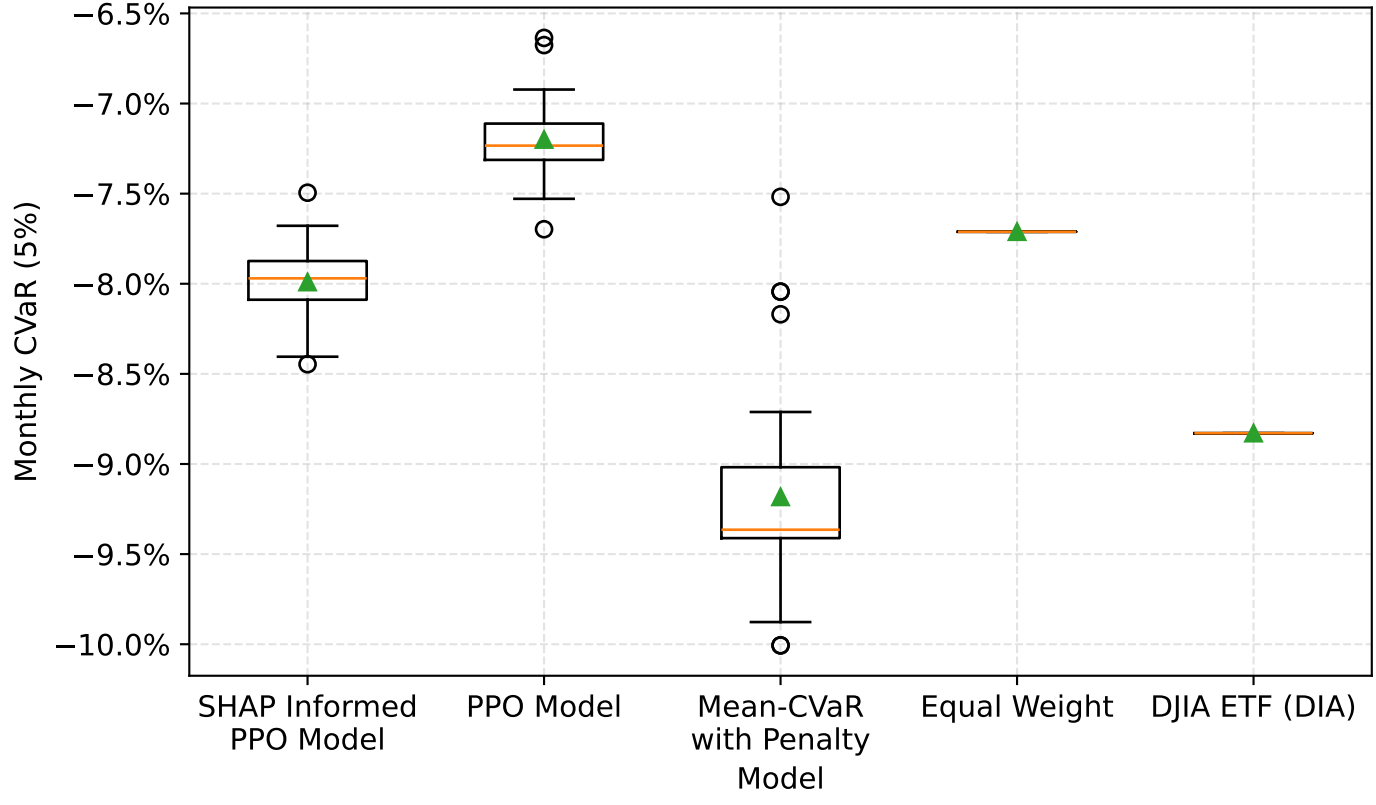

Boxes: IQR; center line: median; diamond = mean; whiskers: non-outlier range.

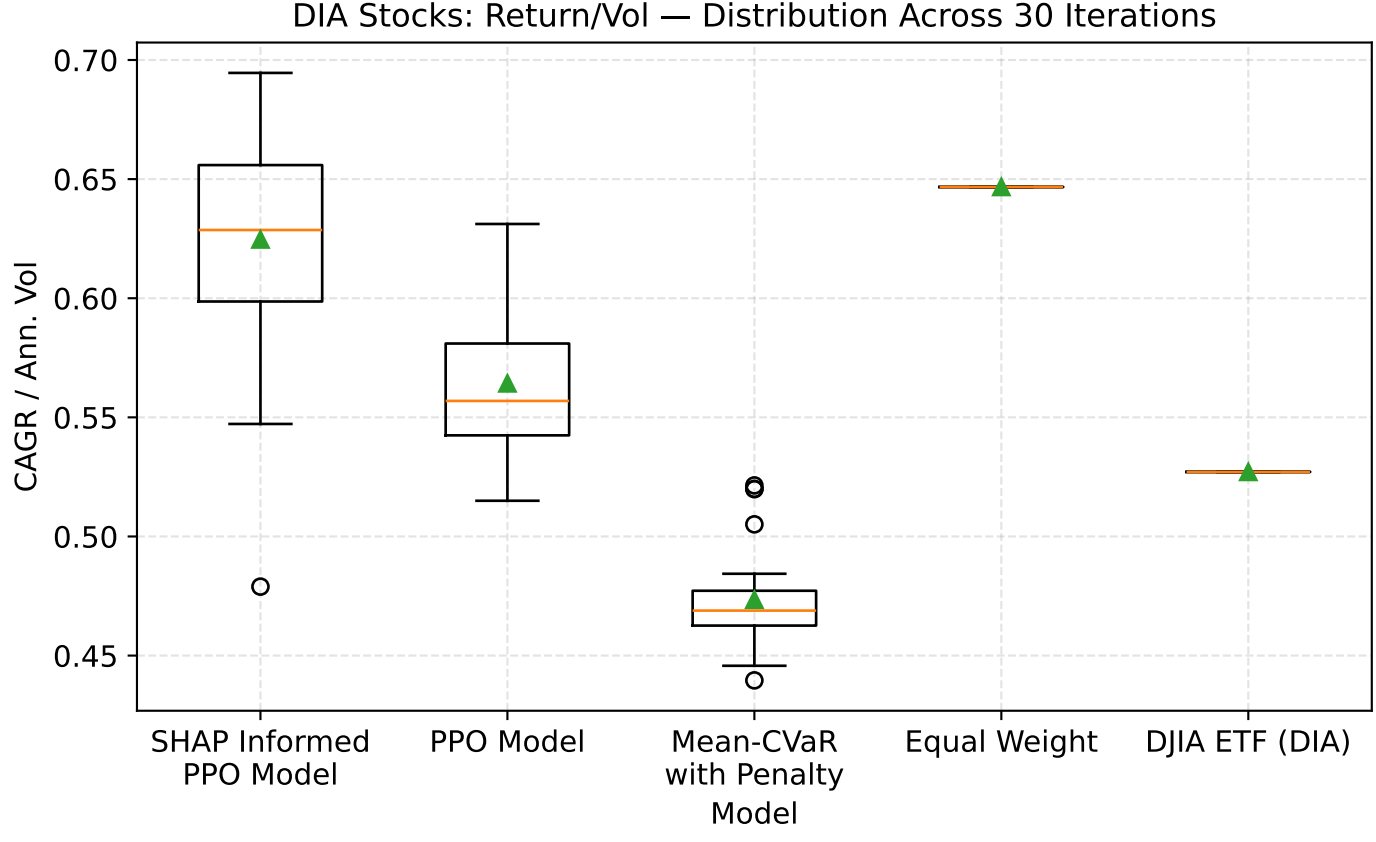

Boxes: IQR; center line: median; diamond = mean; whiskers: non-outlier range.

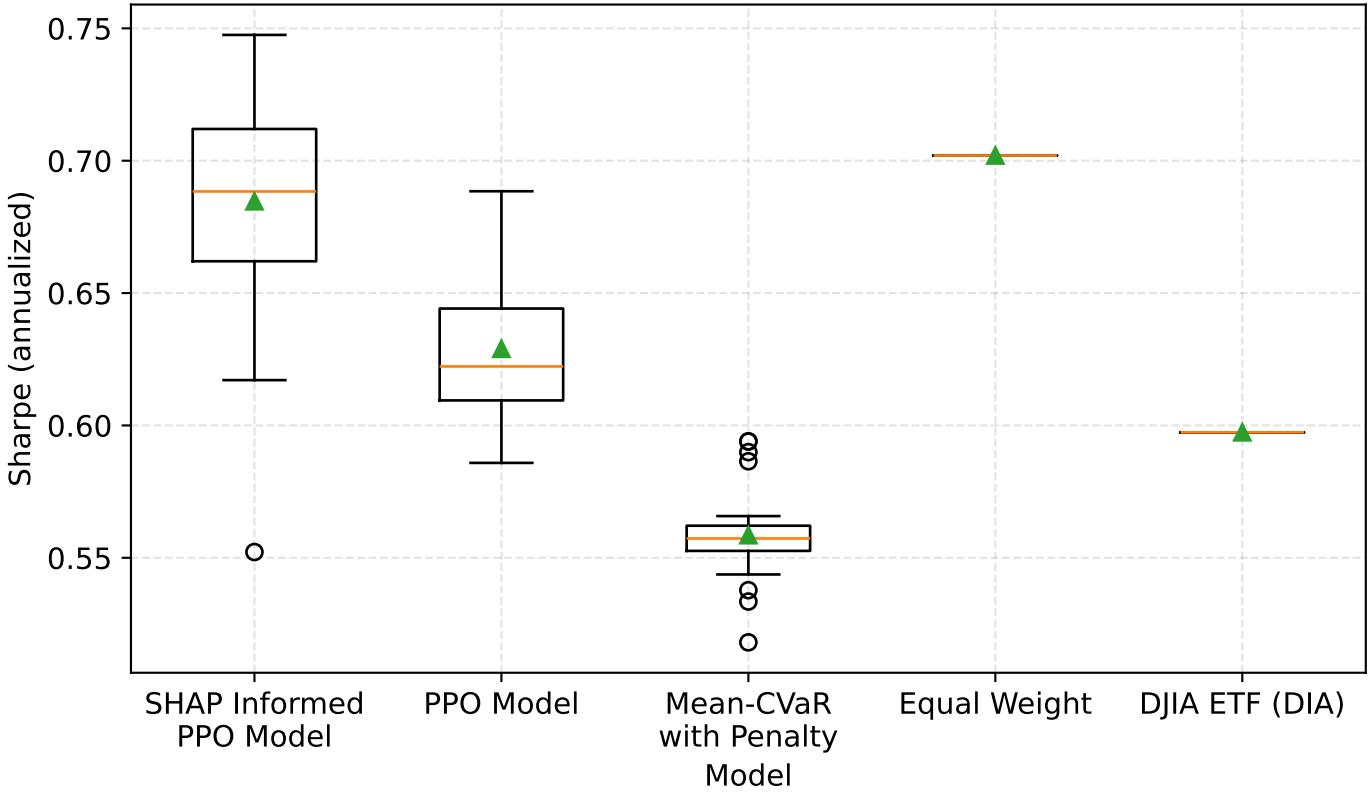

Boxes: IQR; center line: median; diamond = mean; whiskers: non-outlier range.

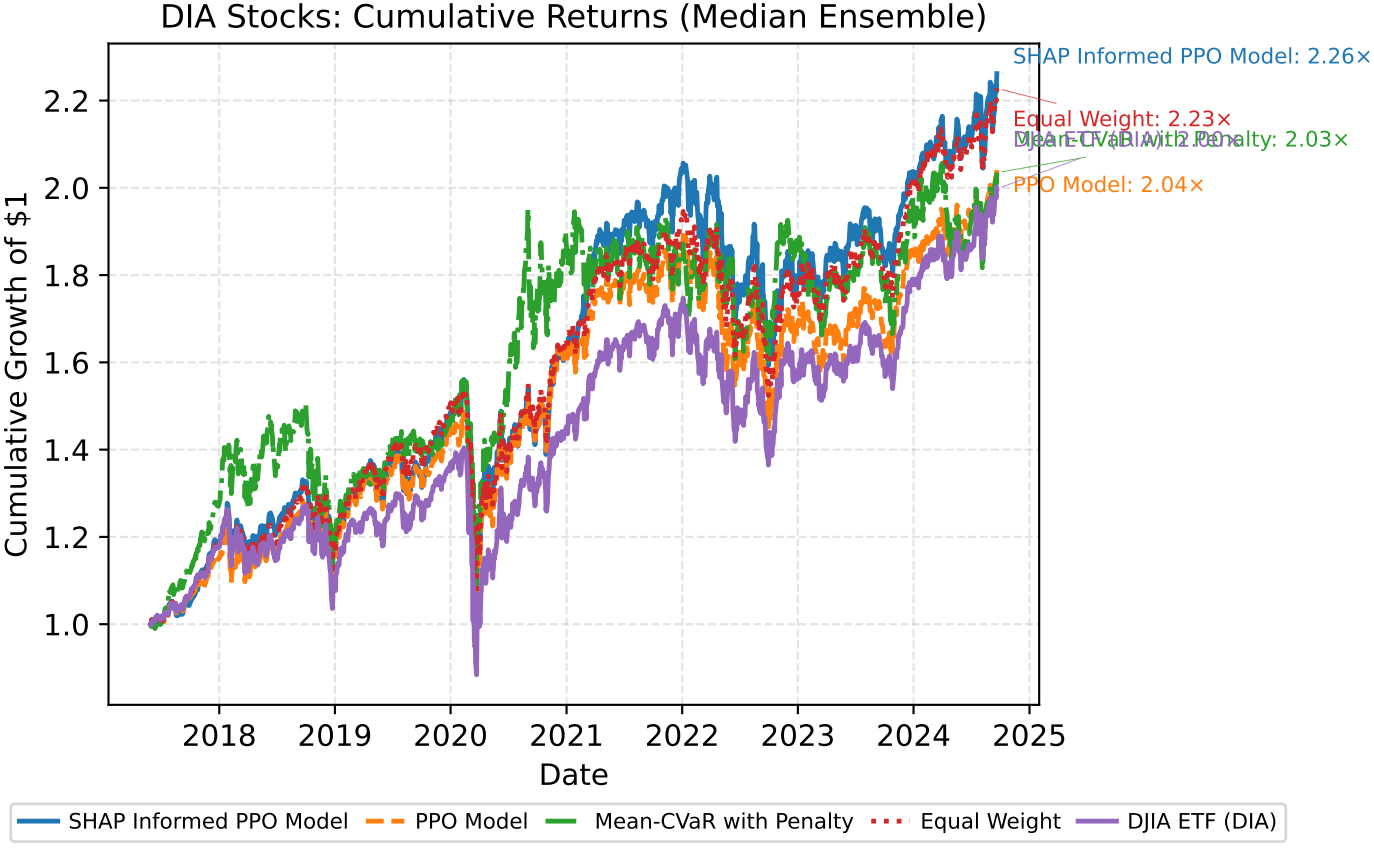

Each series is the median across 30 PPO runs (single path for benchmarks), compounded from daily returns.

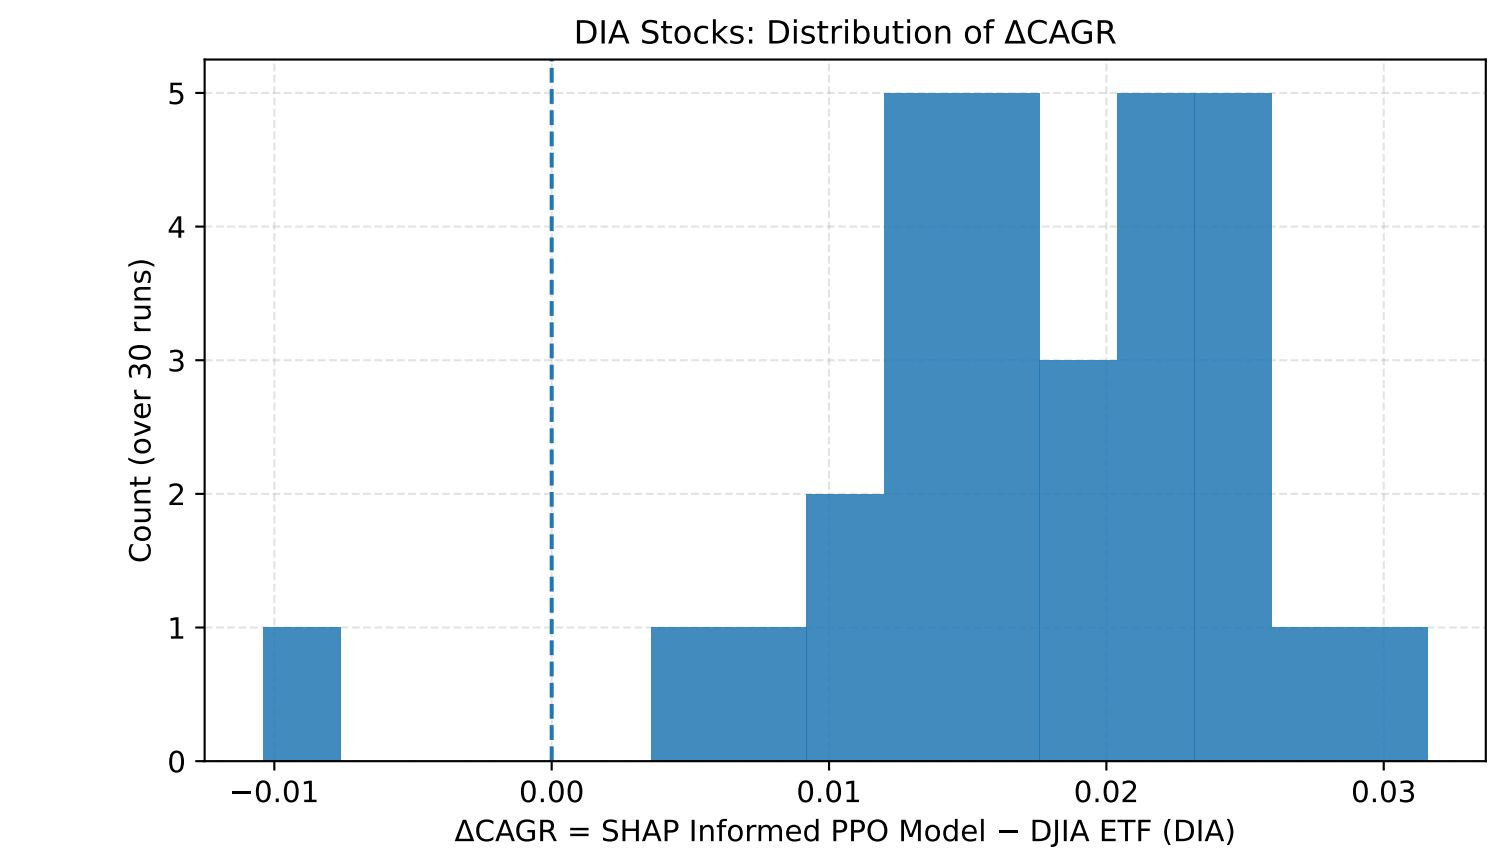

Histogram of per-run metric differences; values > 0 favor SHAP Informed PPO.

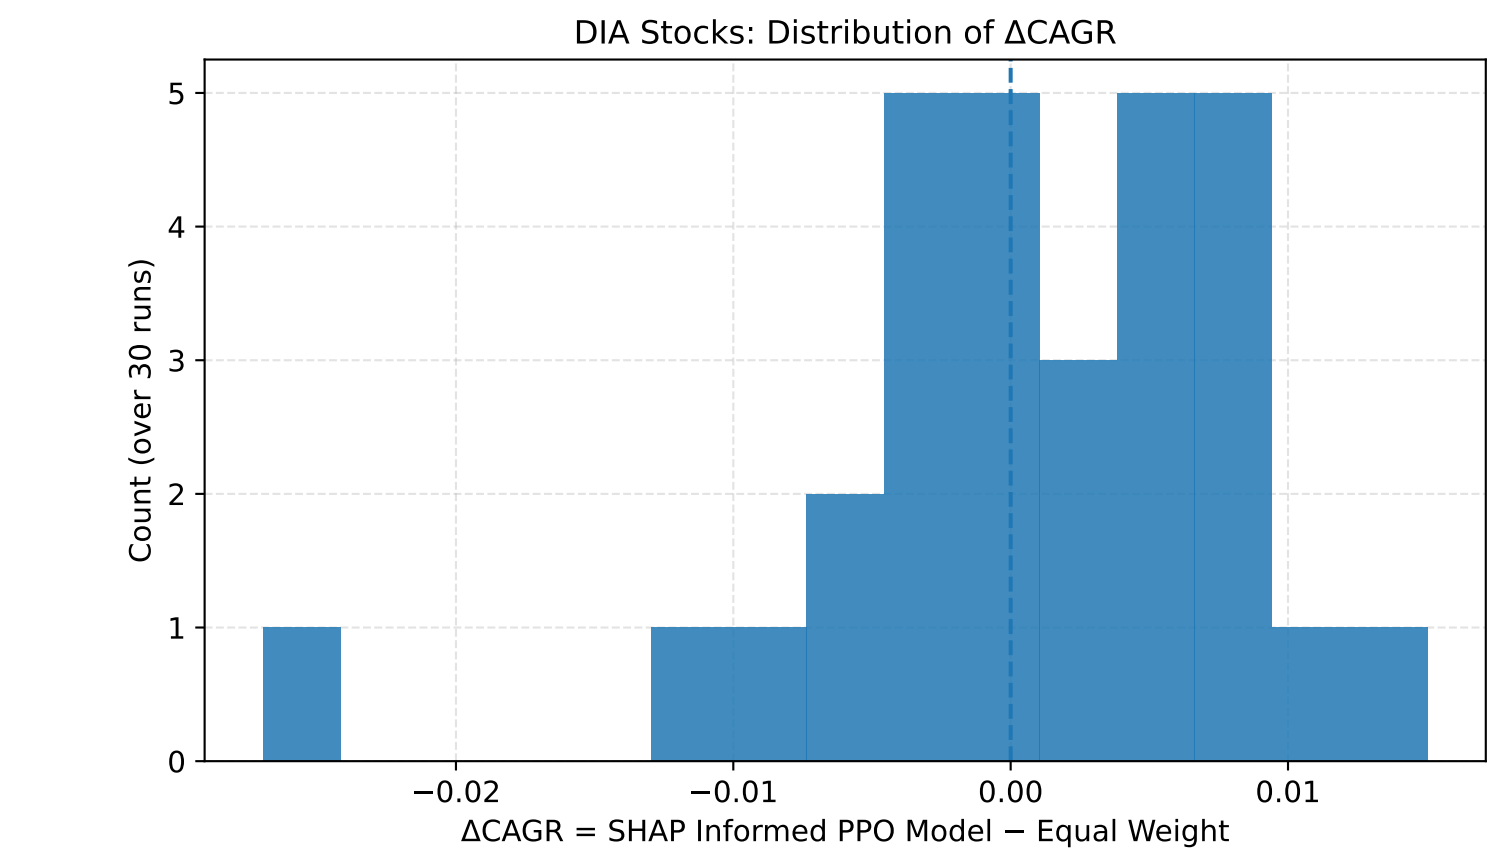

Histogram of per-run metric differences; values > 0 favor SHAP Informed PPO.

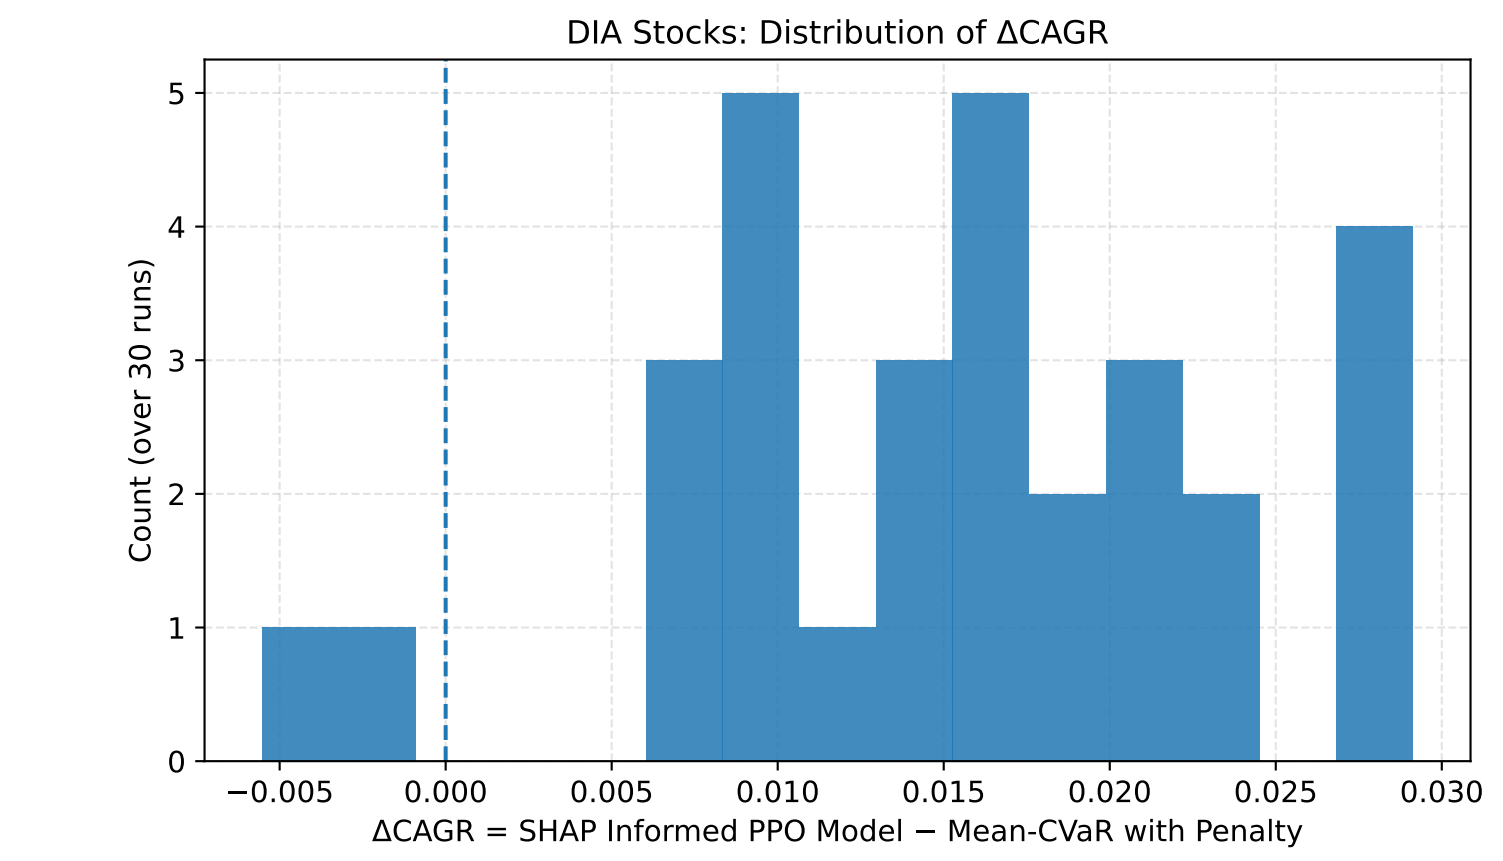

Histogram of per-run metric differences; values > 0 favor SHAP Informed PPO.

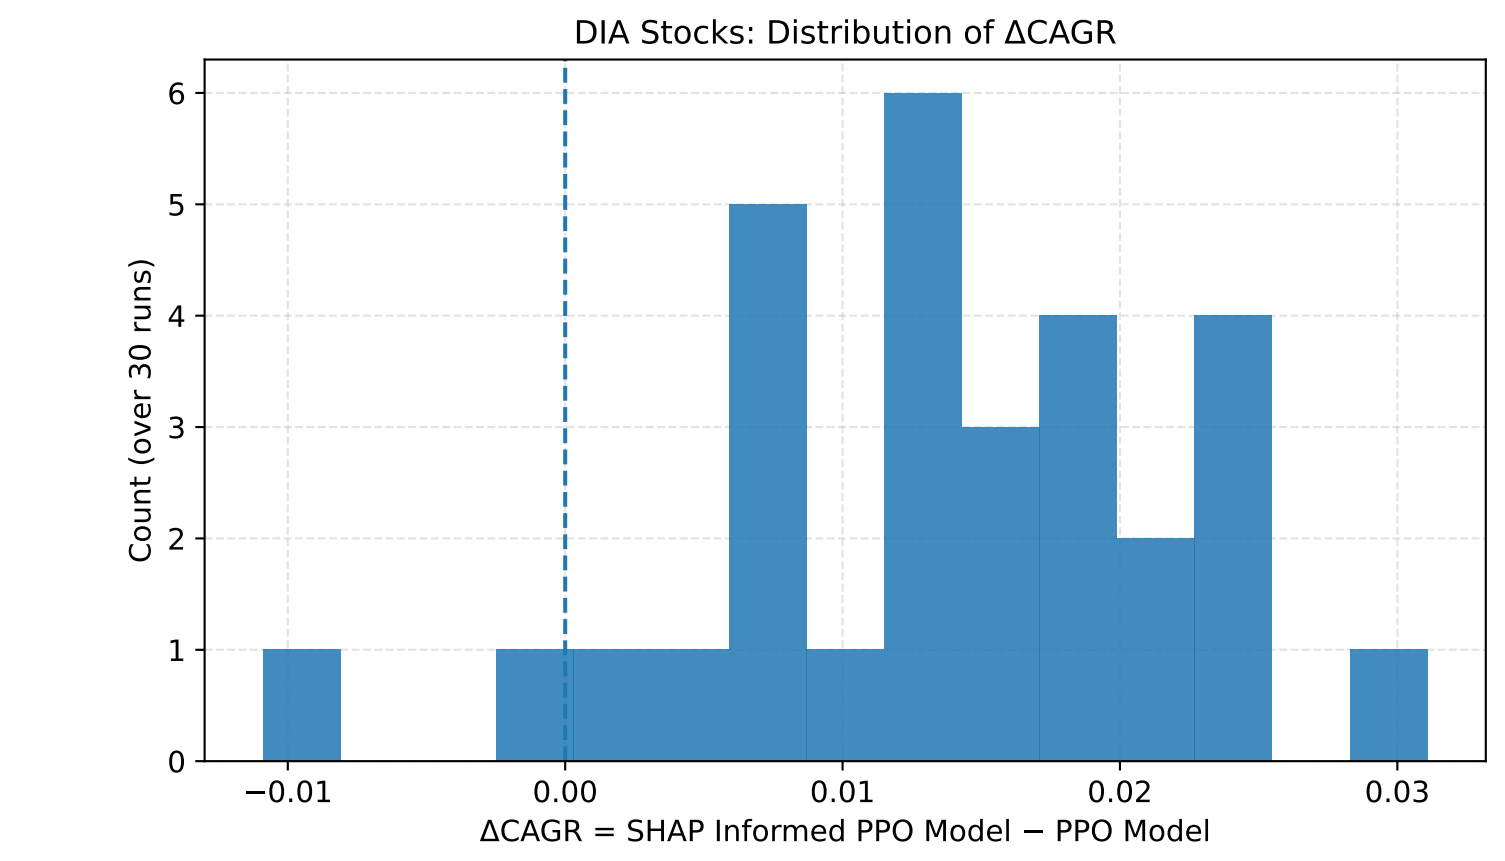

Histogram of per-run metric differences; values > 0 favor SHAP Informed PPO.

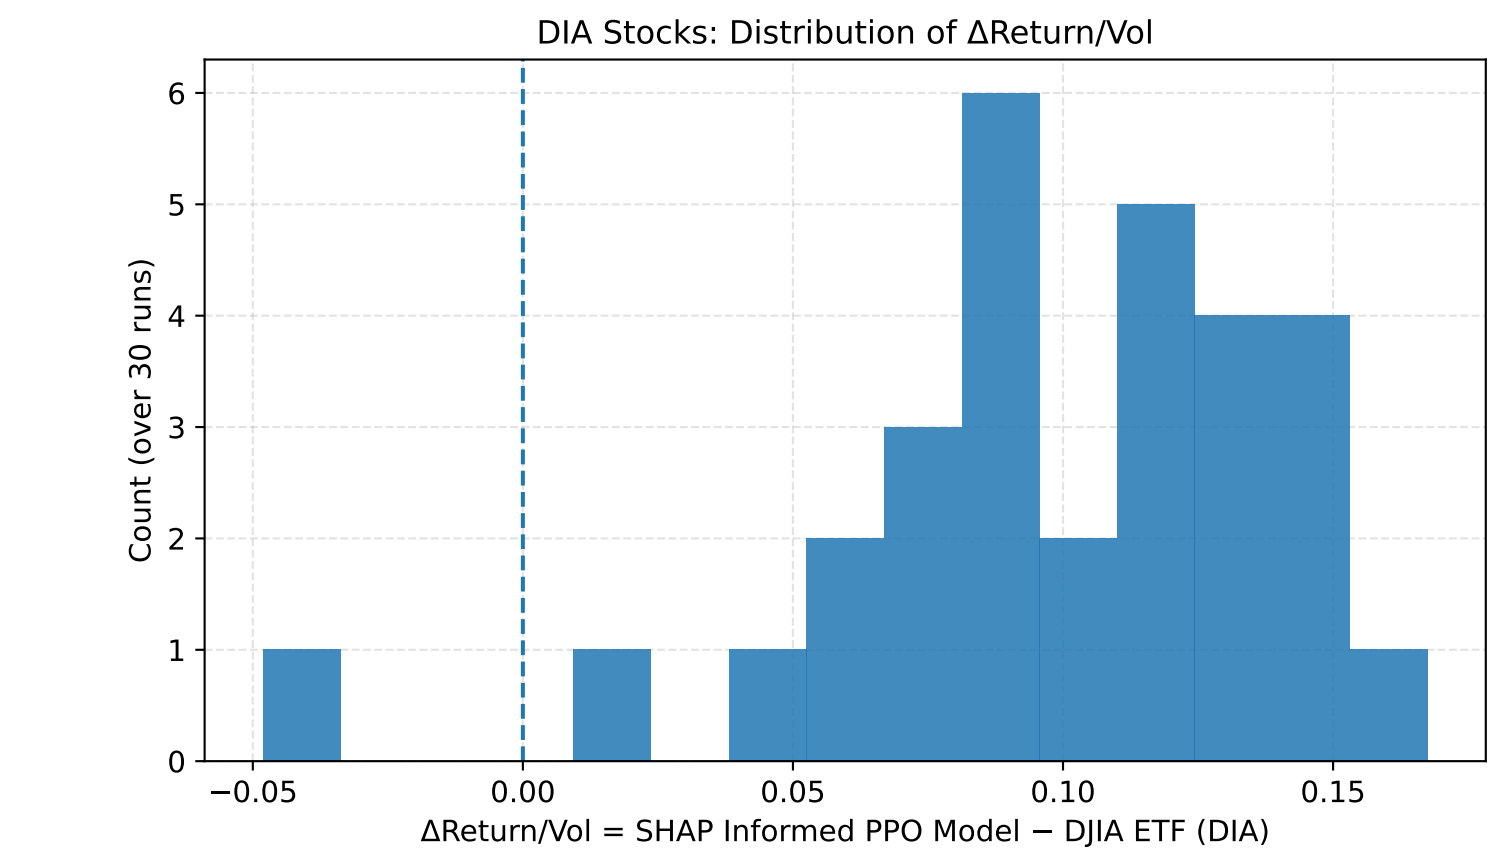

Histogram of per-run metric differences; values > 0 favor SHAP Informed PPO.

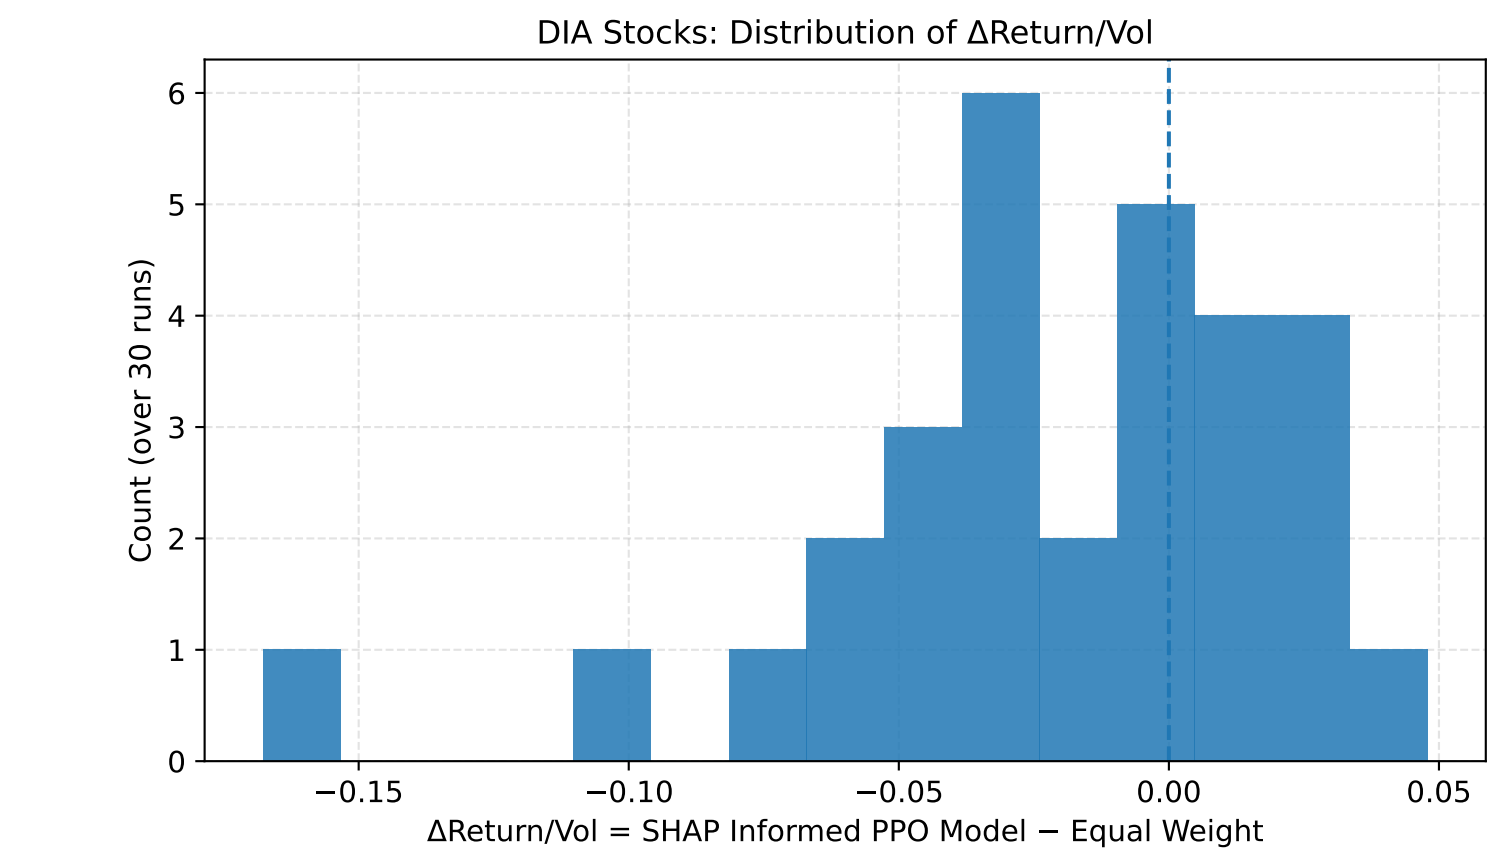

Histogram of per-run metric differences; values > 0 favor SHAP Informed PPO.

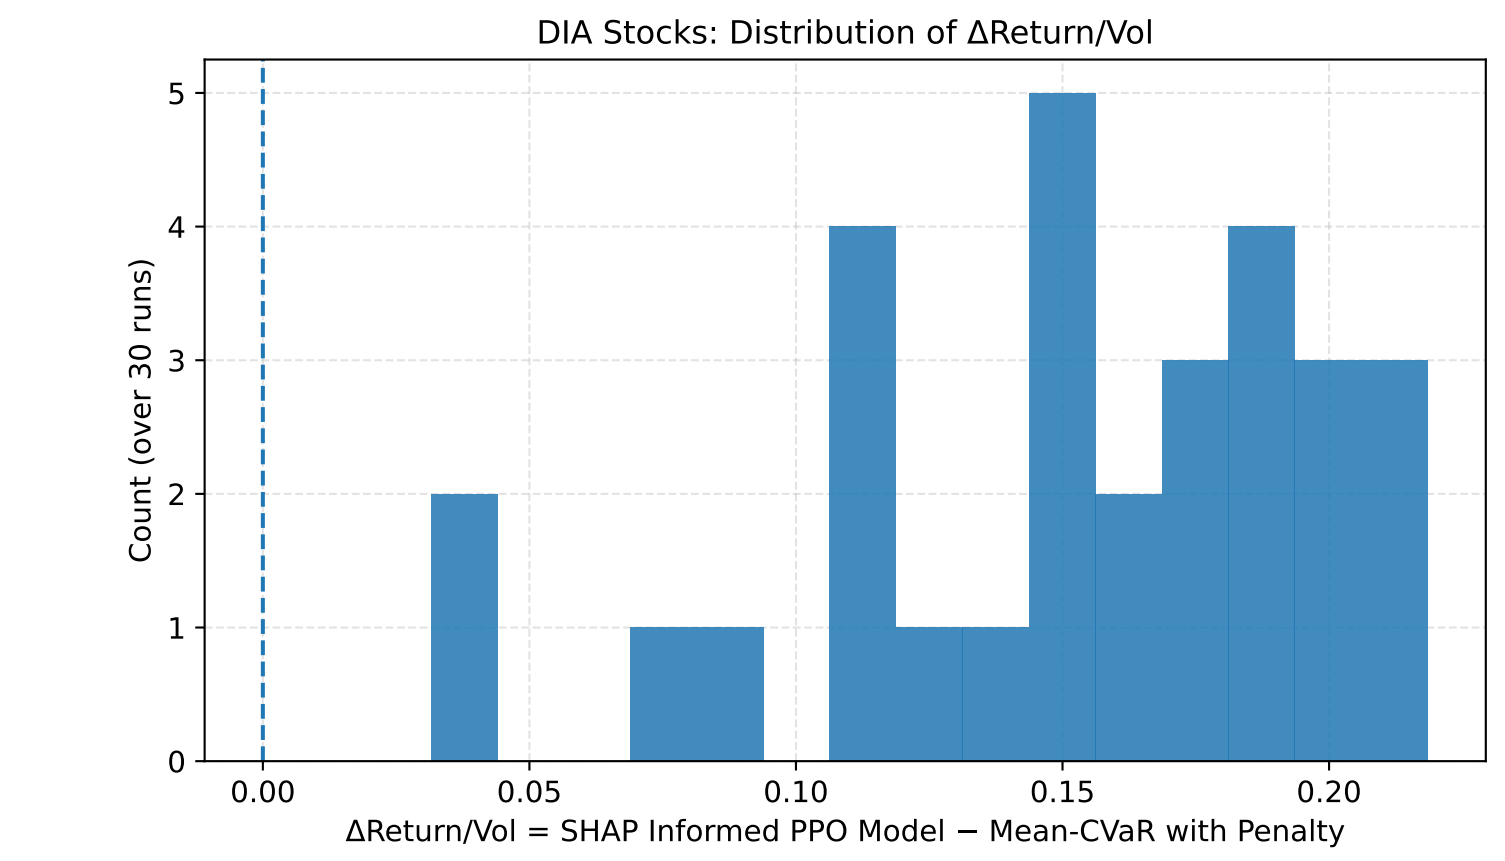

Histogram of per-run metric differences; values > 0 favor SHAP Informed PPO.

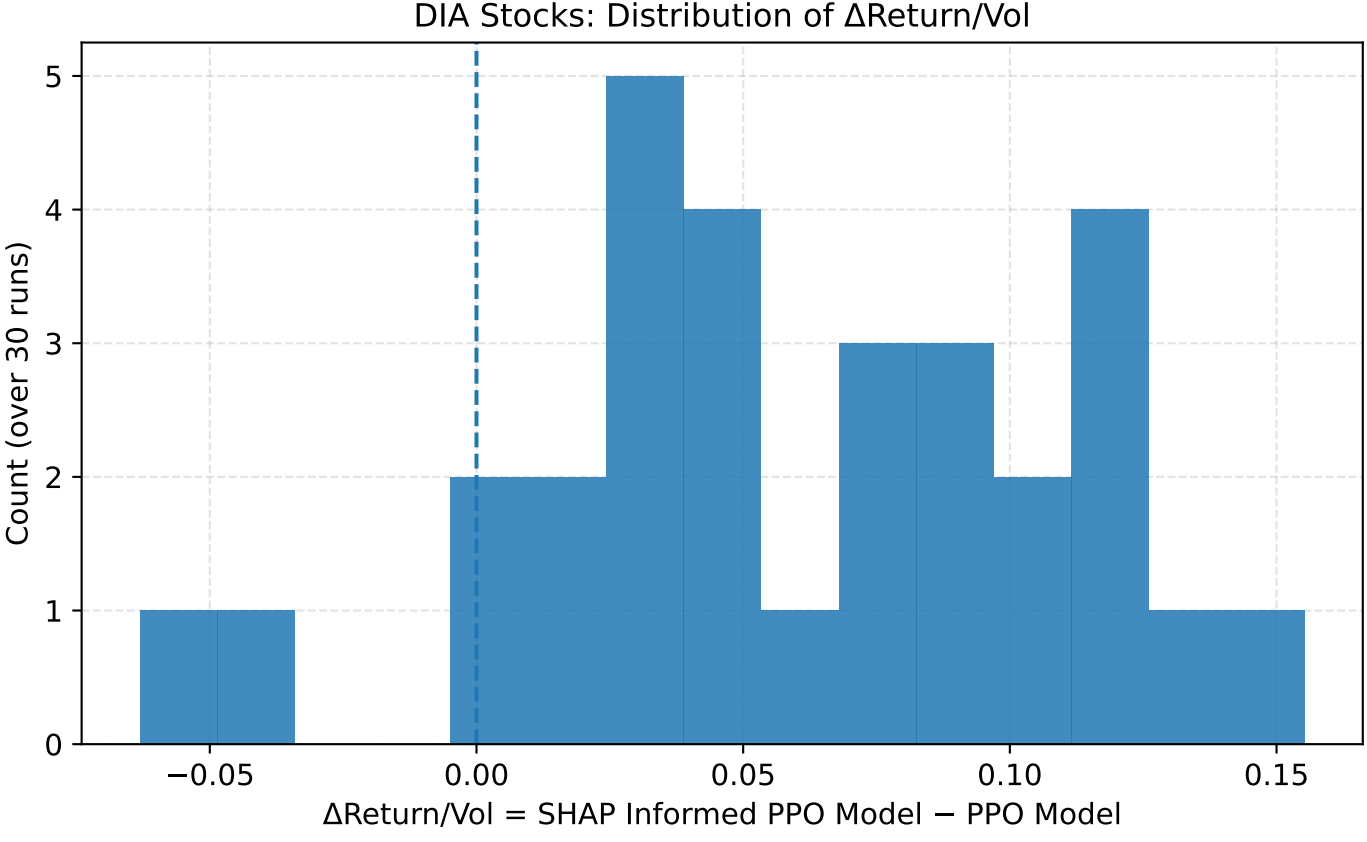

Histogram of per-run metric differences; values > 0 favor SHAP Informed PPO.

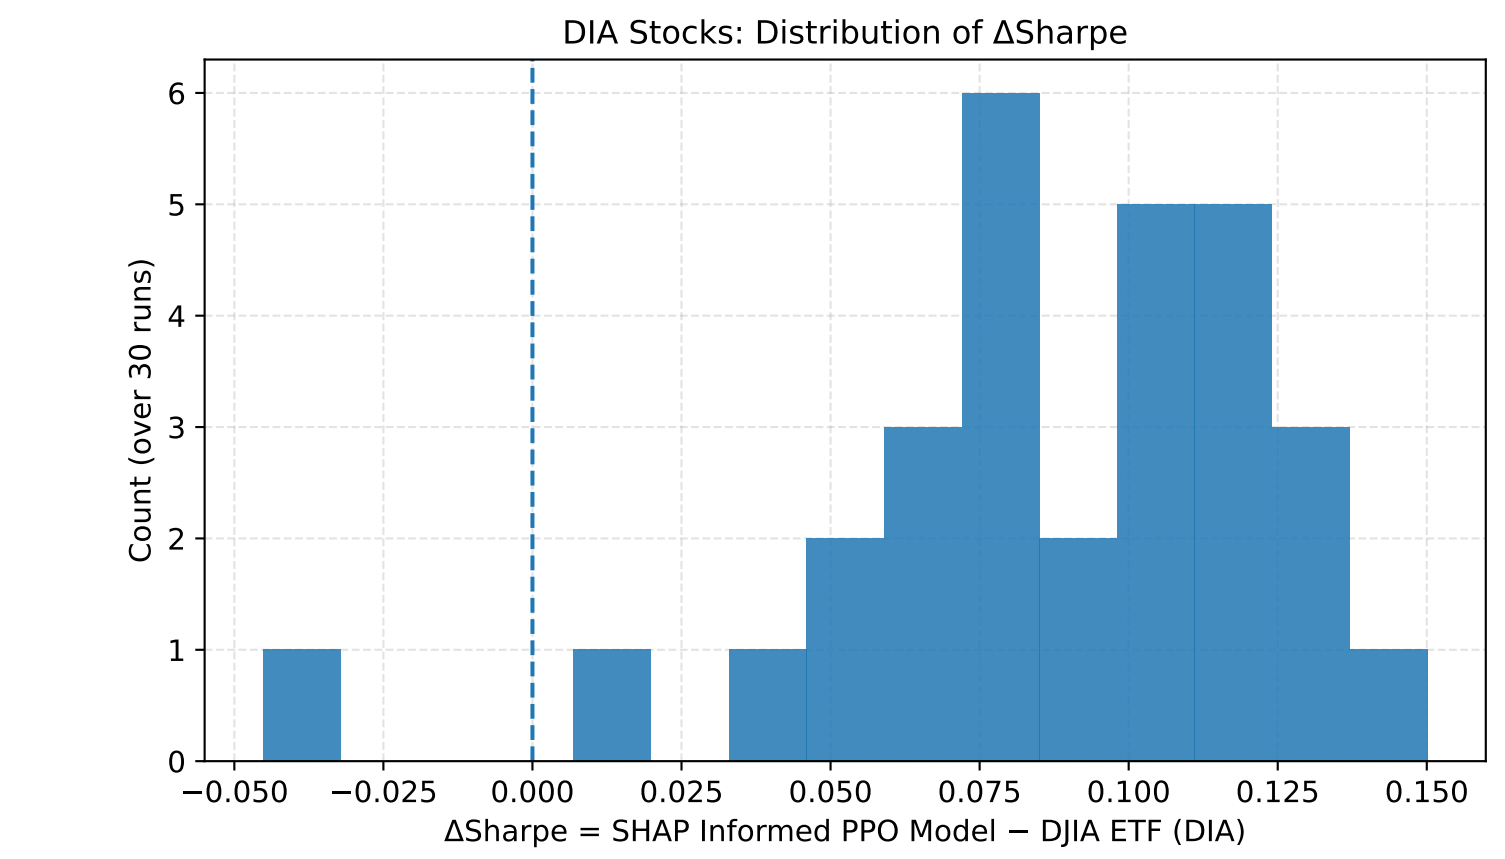

Histogram of per-run metric differences; values > 0 favor SHAP Informed PPO.

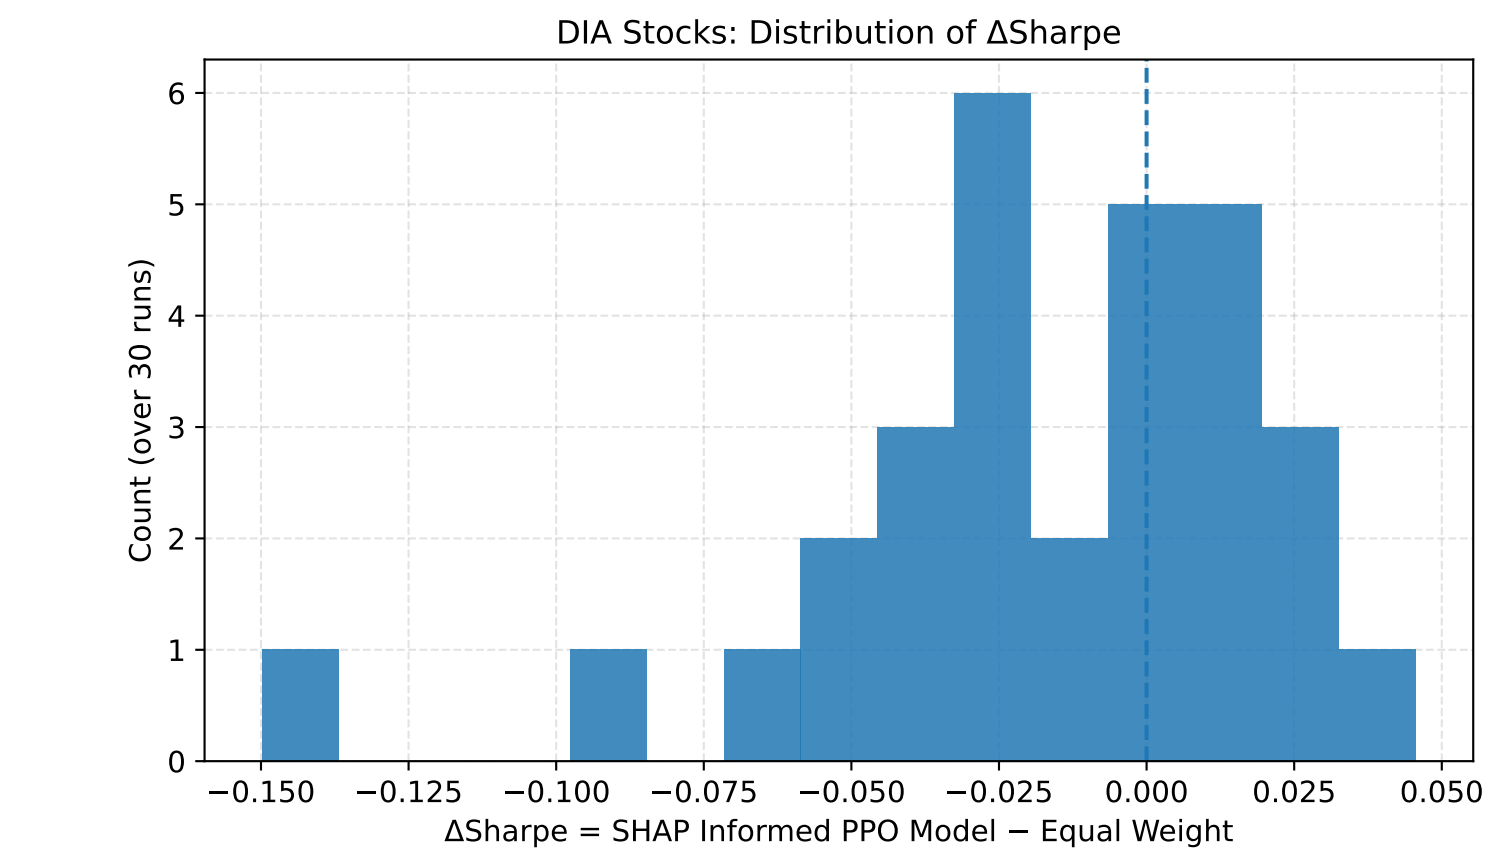

Histogram of per-run metric differences; values > 0 favor SHAP Informed PPO.

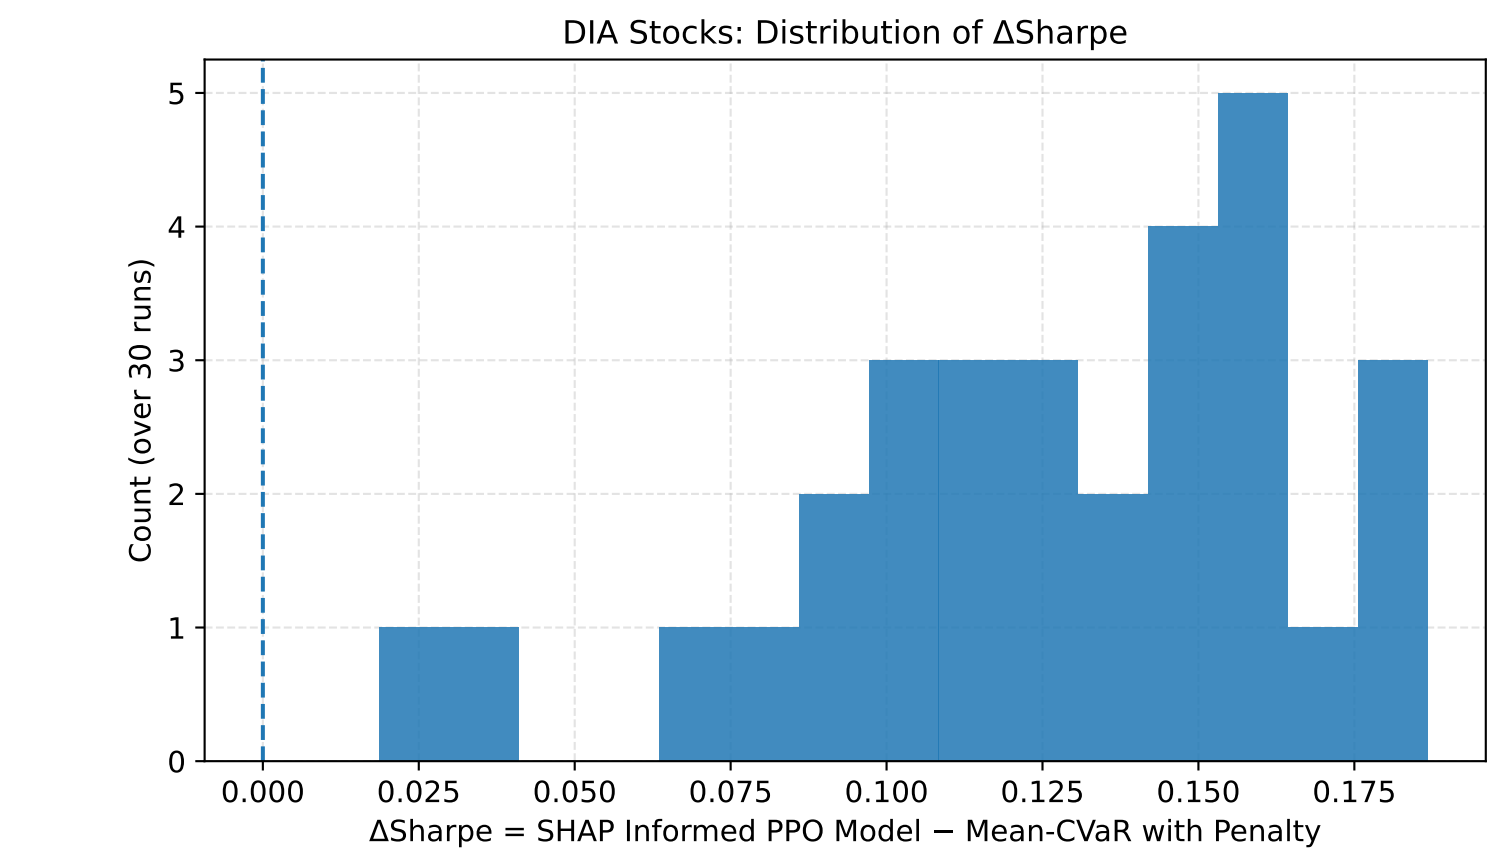

Histogram of per-run metric differences; values > 0 favor SHAP Informed PPO.

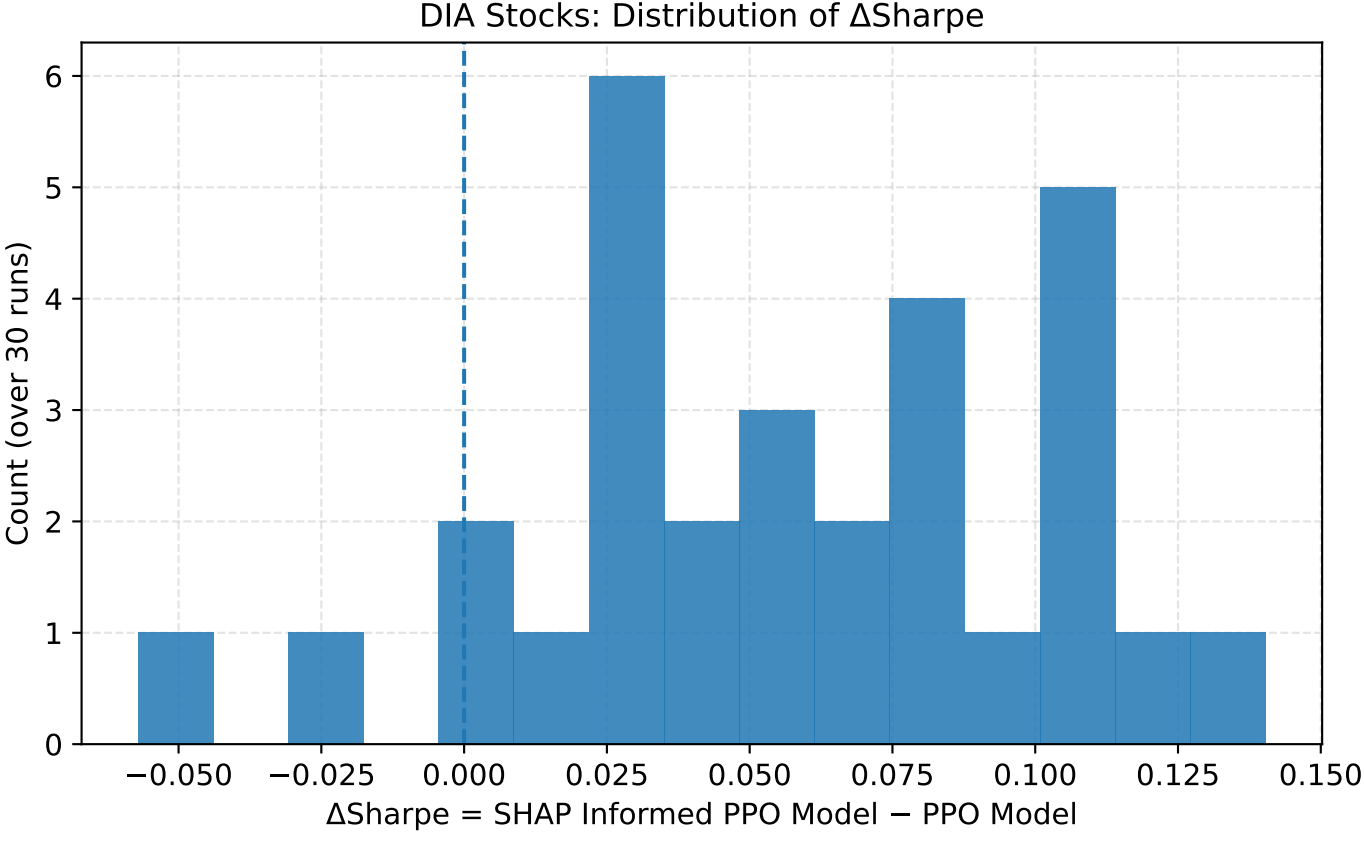

Histogram of per-run metric differences; values > 0 favor SHAP Informed PPO.

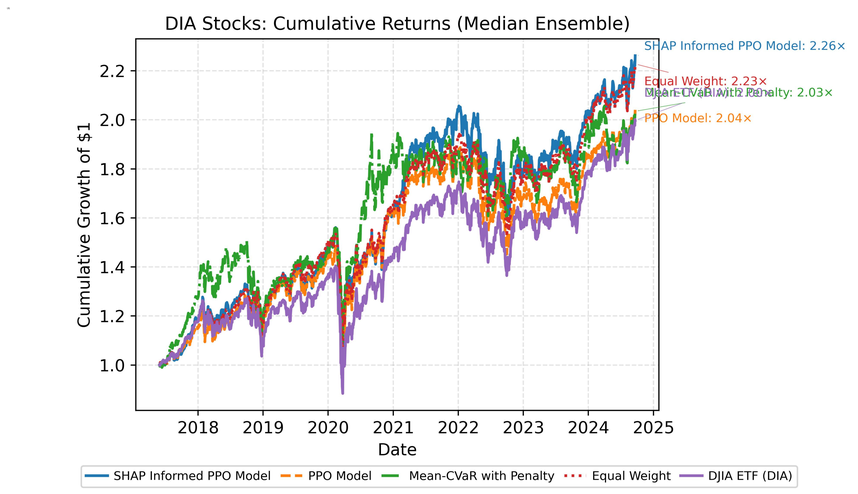

Each series is the median across 30 PPO runs (single path for benchmarks), compounded from daily returns.

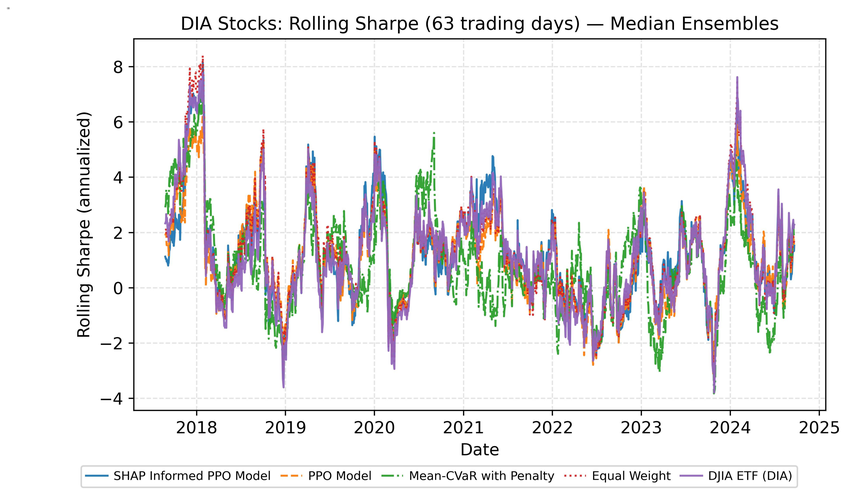

Sharpe computed from rolling mean and std of daily returns (annualized).

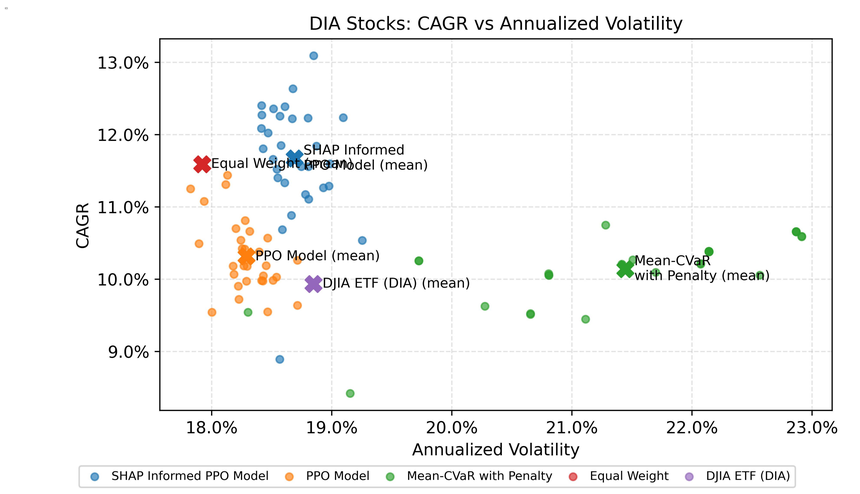

Each point is one PPO run; 'X' marks the model mean across 30 runs.

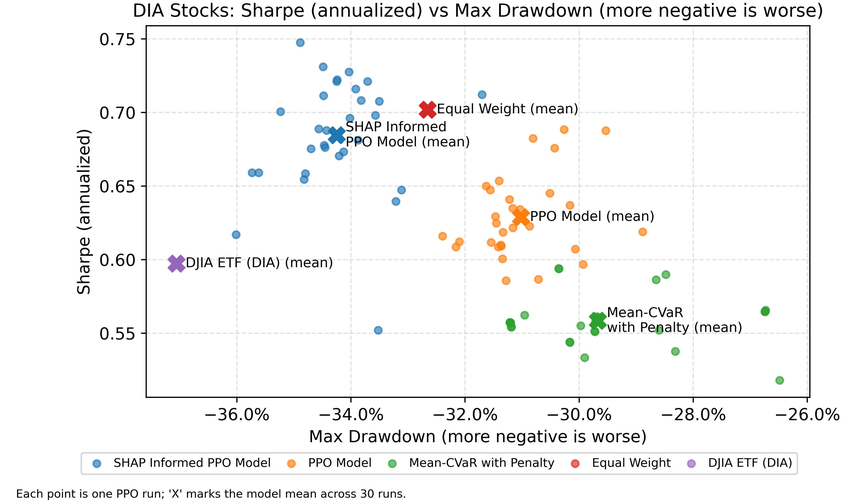

Each point is one PPO run; 'X' marks the model mean across 30 runs.

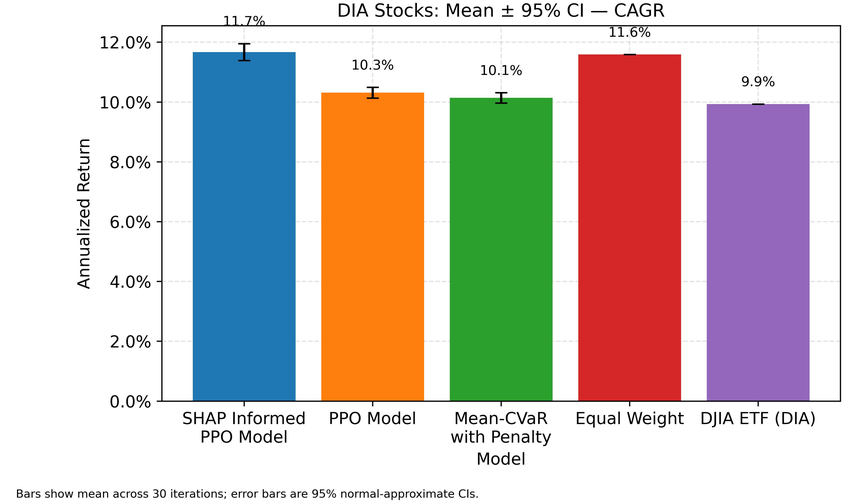

Bars show mean across 30 iterations; error bars are 95% normal-approximate CIs.

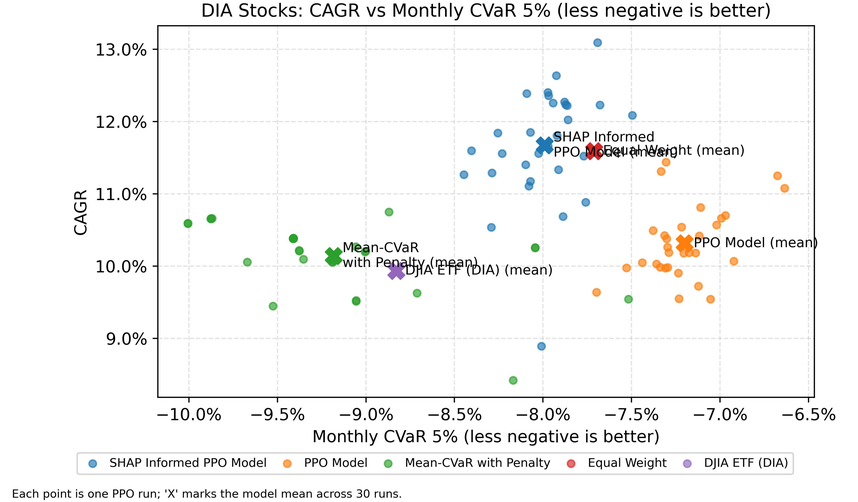

Each point is one PPO run; 'X' marks the model mean across 30 runs.

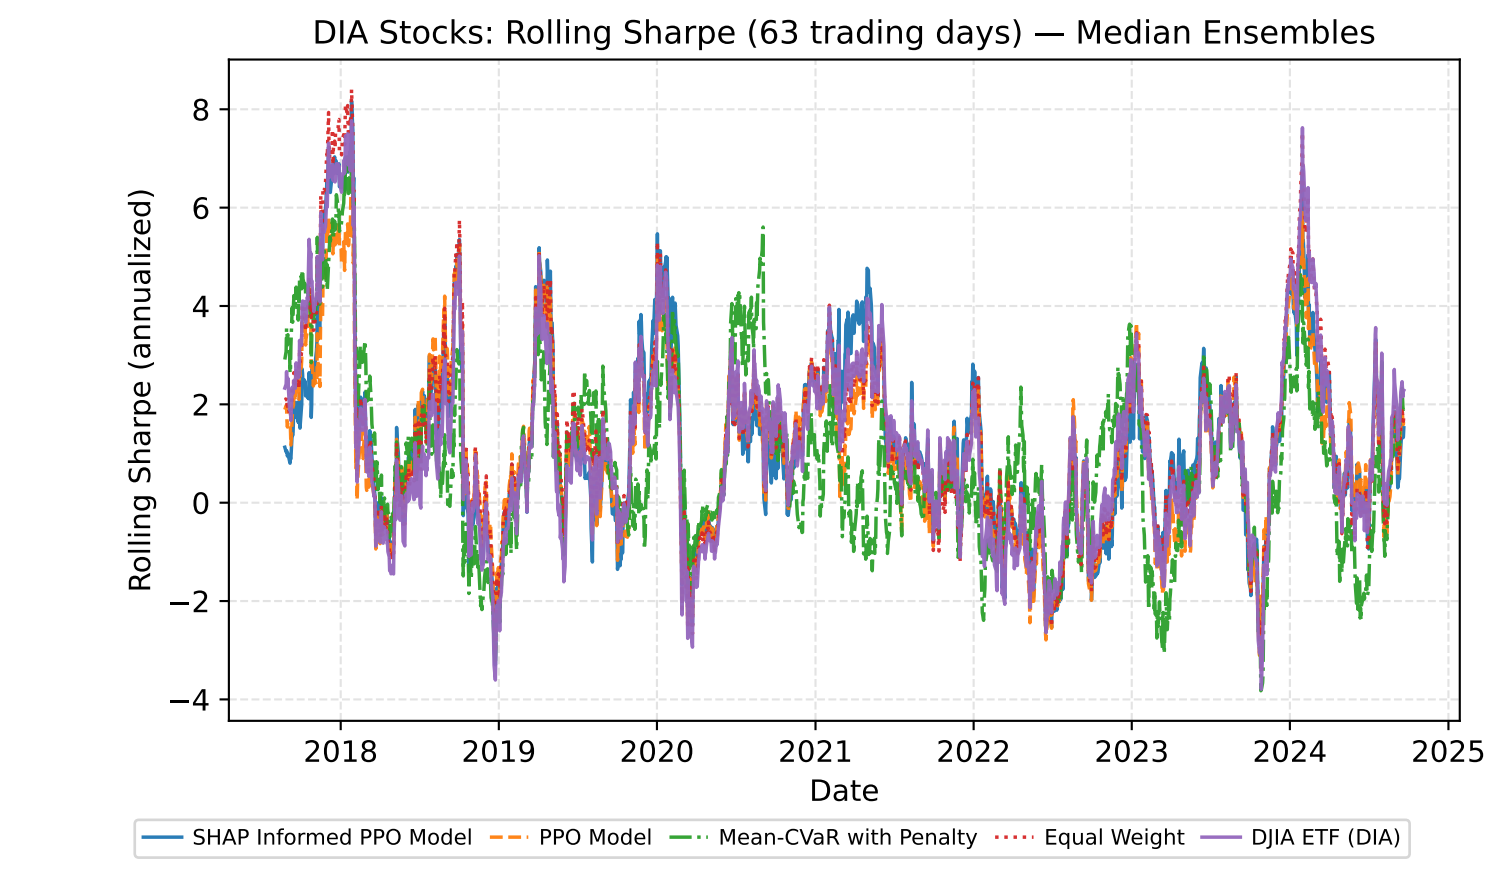

Sharpe computed from rolling mean and std of daily returns (annualized).

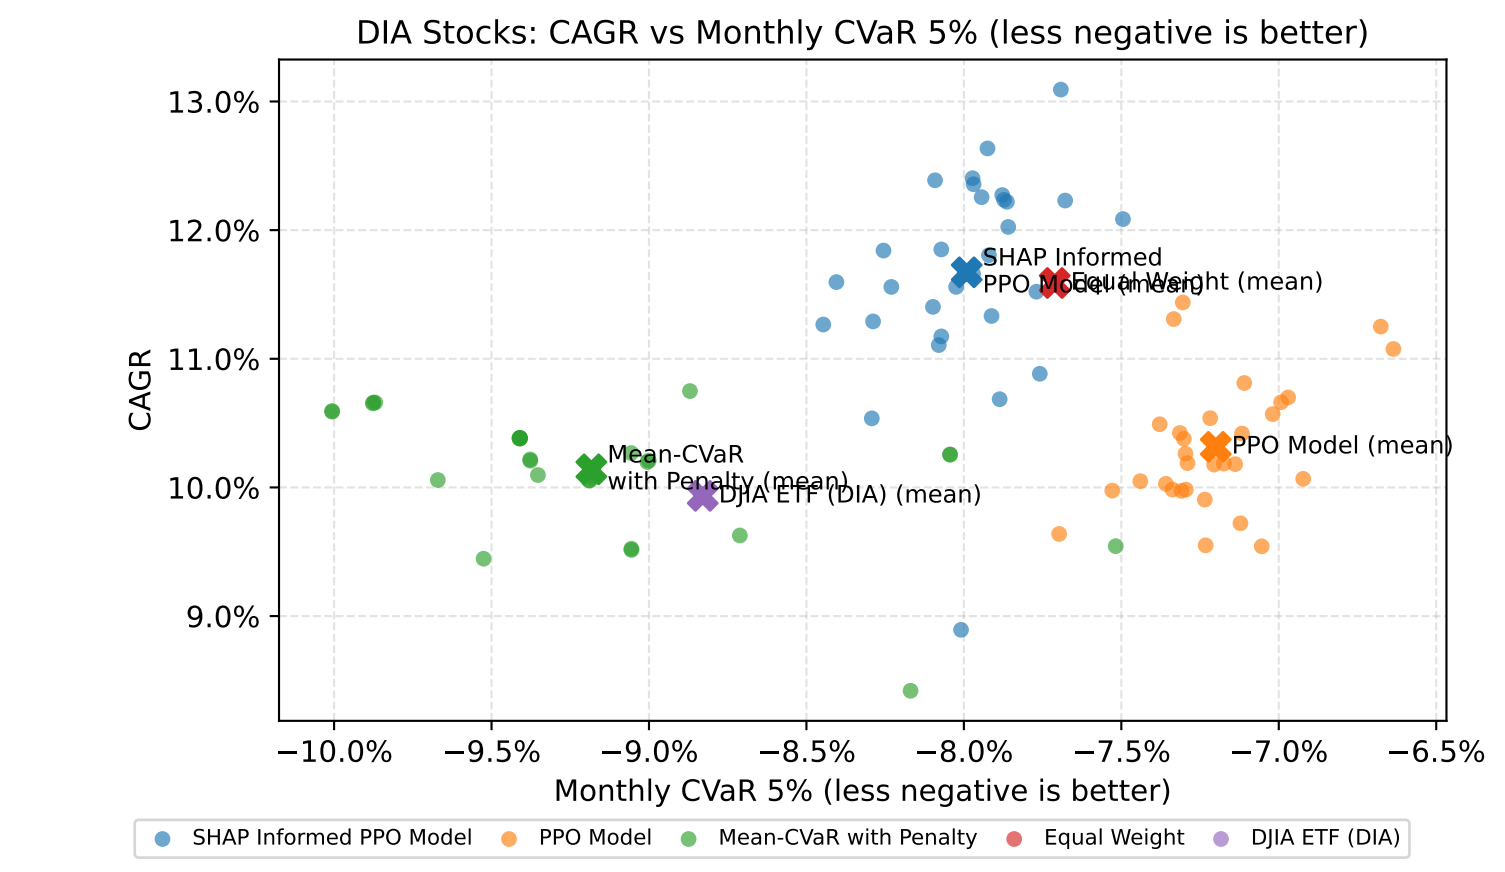

Each point is one PPO run; 'X' marks the model mean across 30 runs.

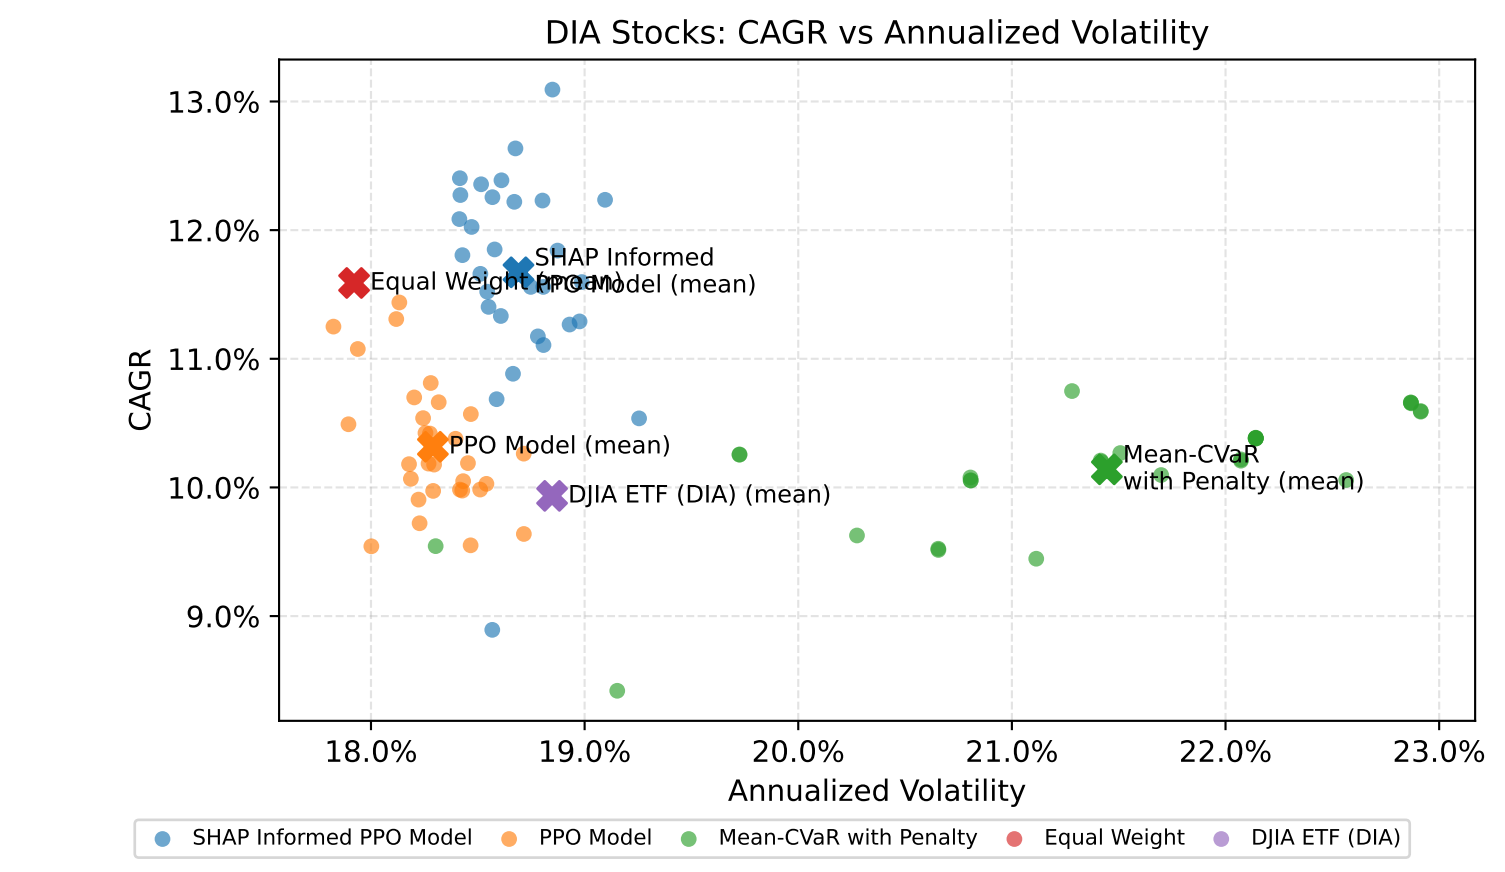

Each point is one PPO run; 'X' marks the model mean across 30 runs.

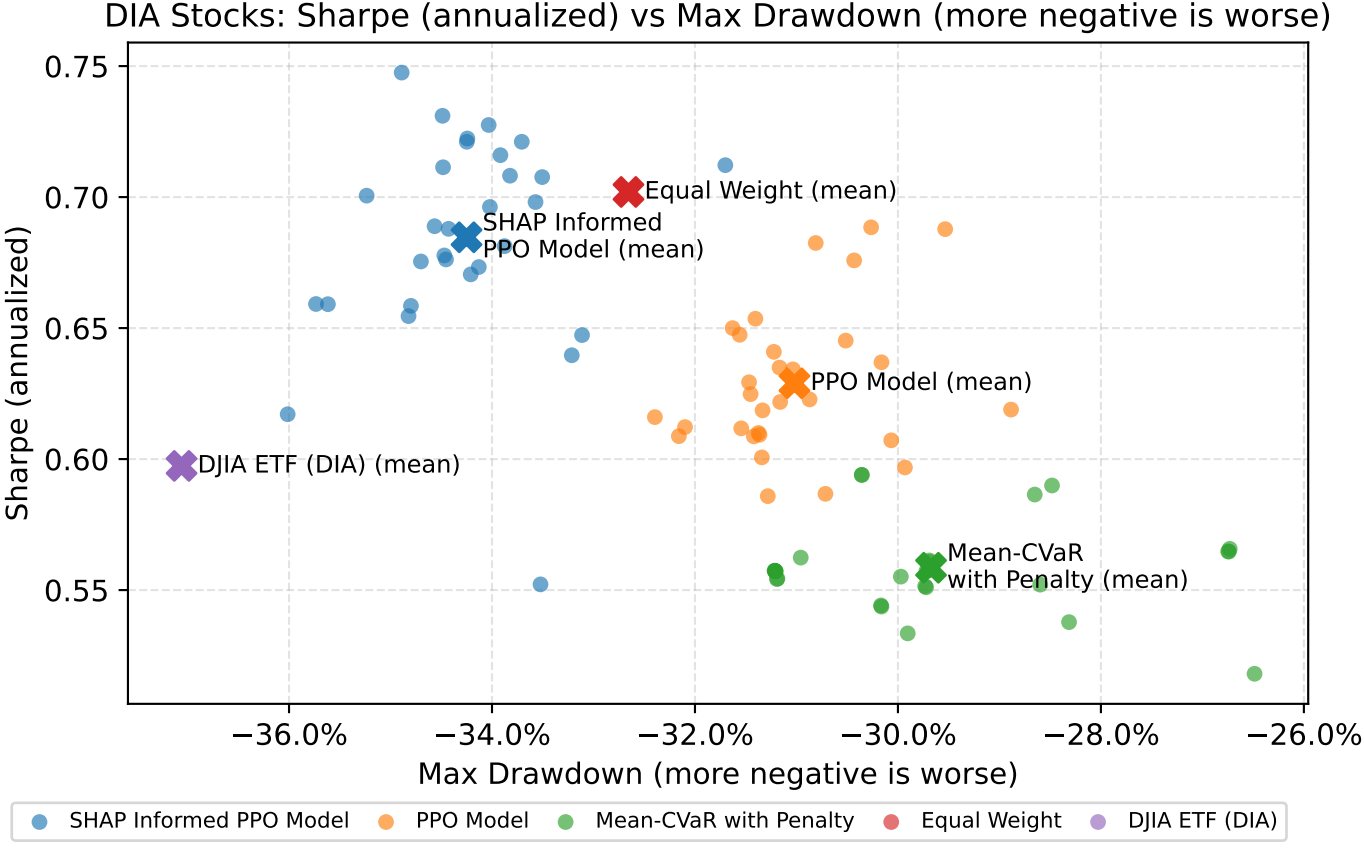

Each point is one PPO run; 'X' marks the model mean across 30 runs.